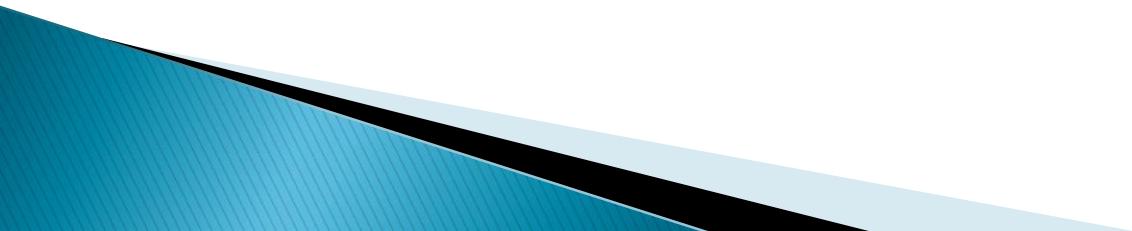


Computer vision tasks



Classification

			
mite black widow cockroach tick starfish	container ship lifeboat amphibian fireboat drilling platform	motor scooter go-kart moped bumper car golfcart	leopard jaguar cheetah snow leopard Egyptian cat
			
grille convertible grille pickup beach wagon fire engine	mushroom agaric mushroom jelly fungus gill fungus dead-man's-fingers	cherry dalmatian grape elderberry ffordshire bullterrier currant	Madagascar cat squirrel monkey spider monkey titi indri howler monkey

Multi-label classification



$y = 0 \ 1 \ 0$



$y = 1 \ 1 \ 0$



$y = 1 \ 0 \ 1$

Activation function → Sigmoid

Binary cross-entropy loss

$$L = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^c y_{ij} \ln p(y_{ij} | x_i) + (1 - y_{ij}) \ln(1 - p(y_{ij} | x_i))$$

Multi-class classification



$y = 1 \ 0 \ 0$



$y = 0 \ 1 \ 0$



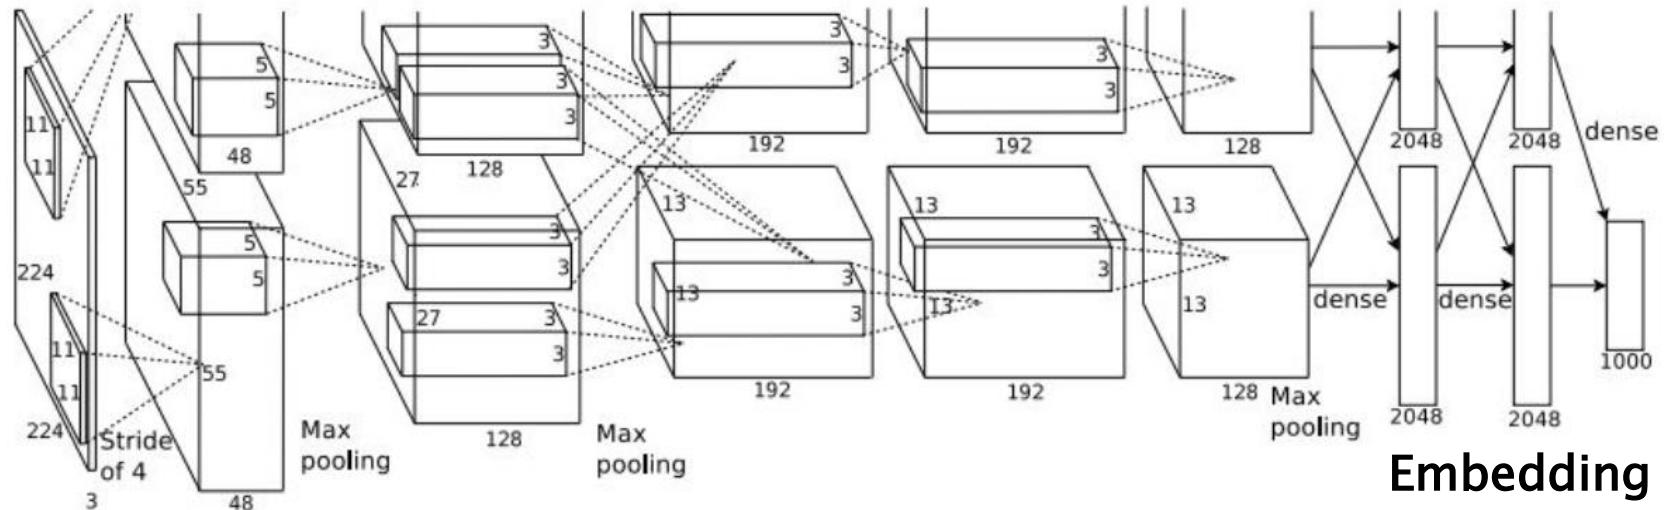
$y = 0 \ 0 \ 1$

Activation function → Softmax

Cross-Entropy Loss Function

$$L = -\frac{1}{N} \sum_i \ln p(c = y_i | x_i)$$

Classification



Embedding

airplane



automobile



bird



cat



deer



dog



Image retrieval

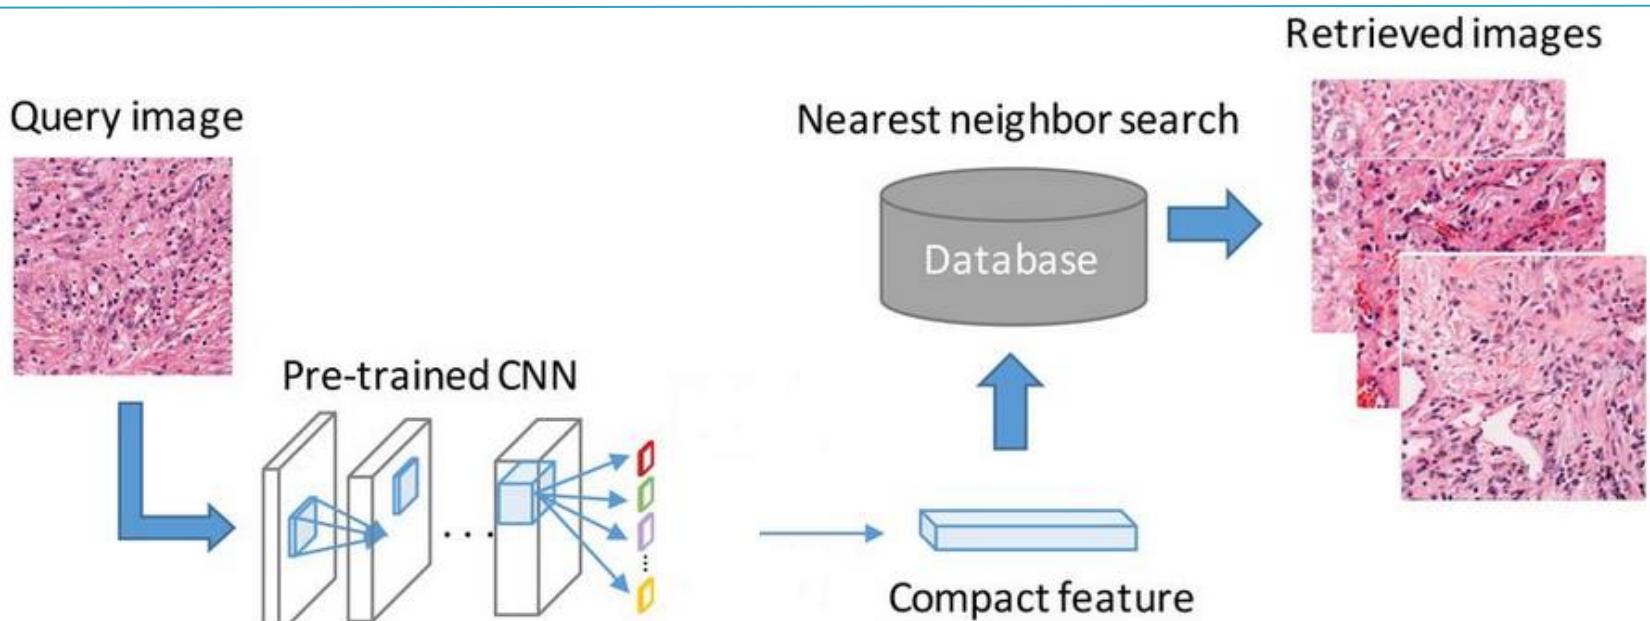
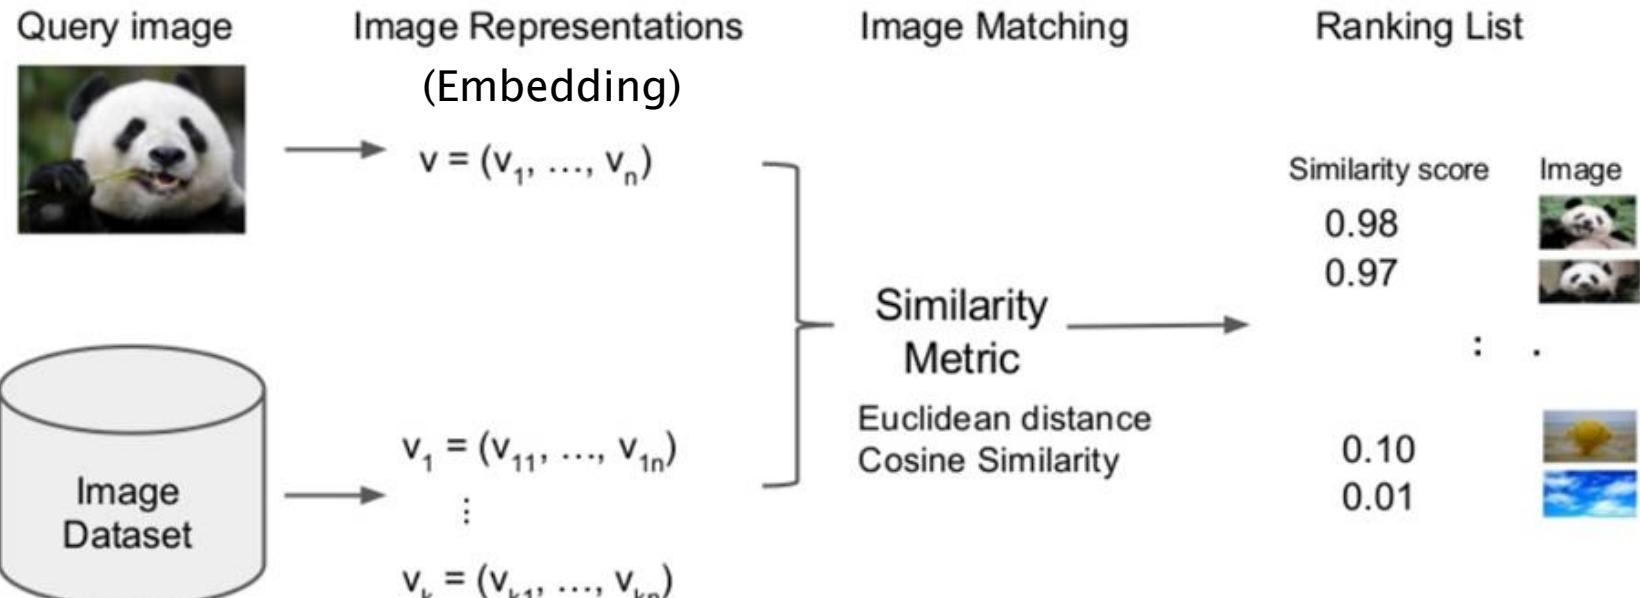


Image tagging

Image Tagging

Выберите изображение

Выберите файл 152804744...1363.jpg



food

Тегировать по категории

Тегировать Расширенное тегирование

Сгенерированные теги

- delicious
- dessert
- eat
- food
- fruit
- fruit_tree
- healthy
- nature
- red
- strawberries

When classification doesn't work

- ▶ extremely unbalanced dataset with a huge number of classes
- ▶ objects of new classes (not from the training set)

Megaface

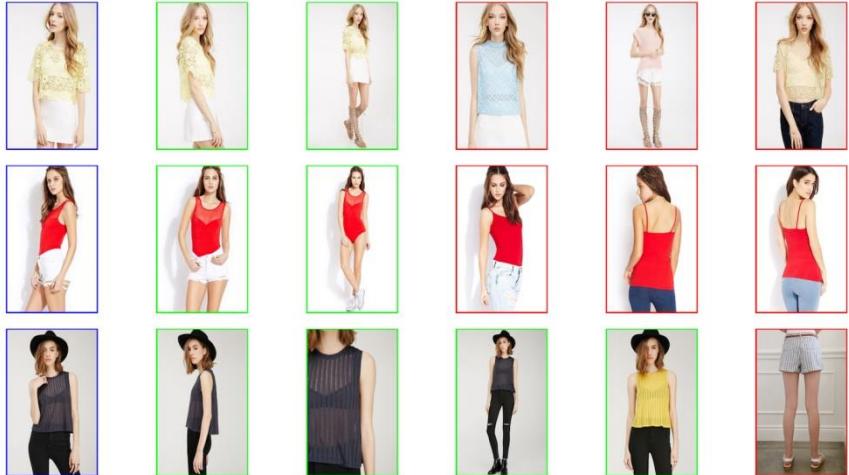


17 categories of clothing (jackets, jeans, shorts, etc.)
~8 thousand classes (articles of specific products)
The median class size is 5 images.

Distractors

1 Million Photos

690,572 Unique Users



Training Set

4.7 Million Photos

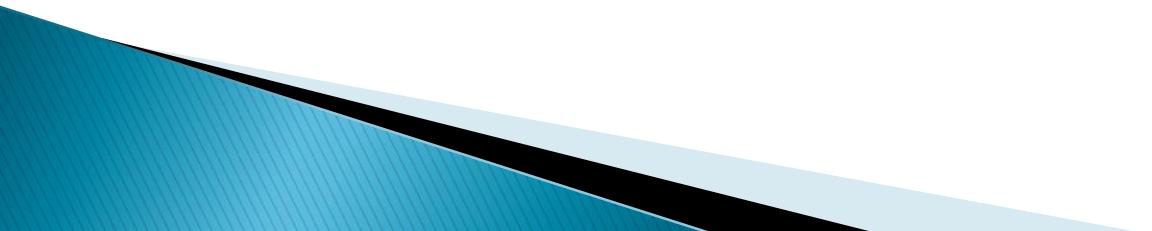
672,057 Unique Identities

7 Mean photos / person (3 min, 2469 max)



DeepFashion

Metric Learning



Triplet Loss



Figure 2. **Model structure.** Our network consists of a batch input layer and a deep CNN followed by L_2 normalization, which results in the face embedding. This is followed by the triplet loss during training.

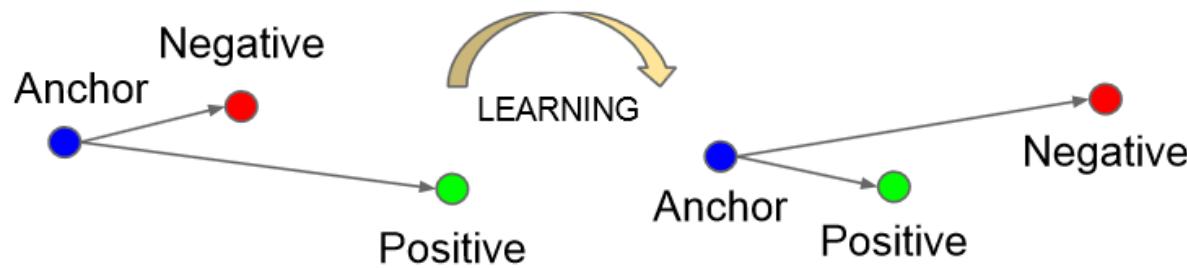


Figure 3. The **Triplet Loss** minimizes the distance between an *anchor* and a *positive*, both of which have the same identity, and maximizes the distance between the *anchor* and a *negative* of a different identity.

$$L = \sum_i^N \left[\|f(x_i^a) - f(x_i^p)\|_2^2 - \|f(x_i^a) - f(x_i^n)\|_2^2 + \alpha \right]_+$$

TensorFlow Addons Losses: TripletSemiHardLoss

[Run in Google Colab](#)[View source on GitHub](#)[Download notebook](#)

Overview

This notebook will demonstrate how to use the `TripletSemiHardLoss` function in TensorFlow Addons.

Resources:

- [FaceNet: A Unified Embedding for Face Recognition and Clustering](#)
- [Oliver Moindrot's blog does an excellent job of describing the algorithm in detail](#)

TripletLoss

As first introduced in the FaceNet paper, TripletLoss is a loss function that trains a neural network to closely embed features of the same class while maximizing the distance between embeddings of different classes. To do this an anchor is chosen along with one negative and one positive sample.

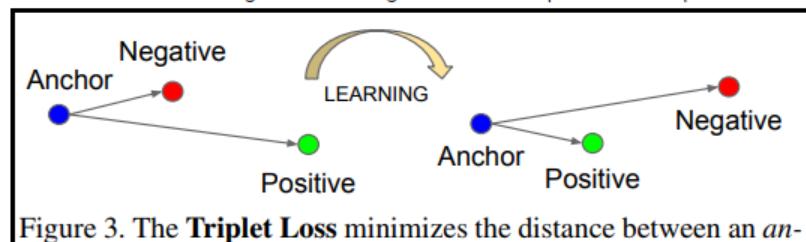


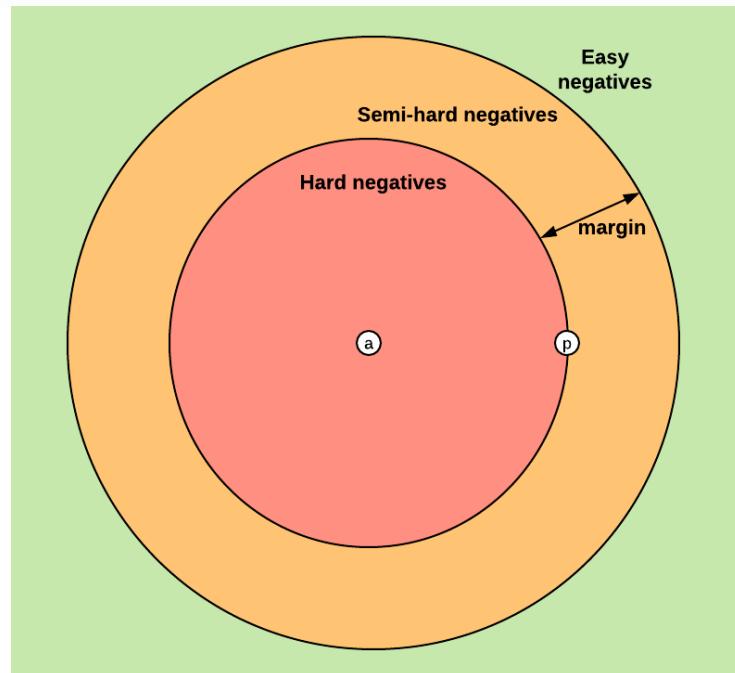
Figure 3. The **Triplet Loss** minimizes the distance between an *an-*

Triplet mining

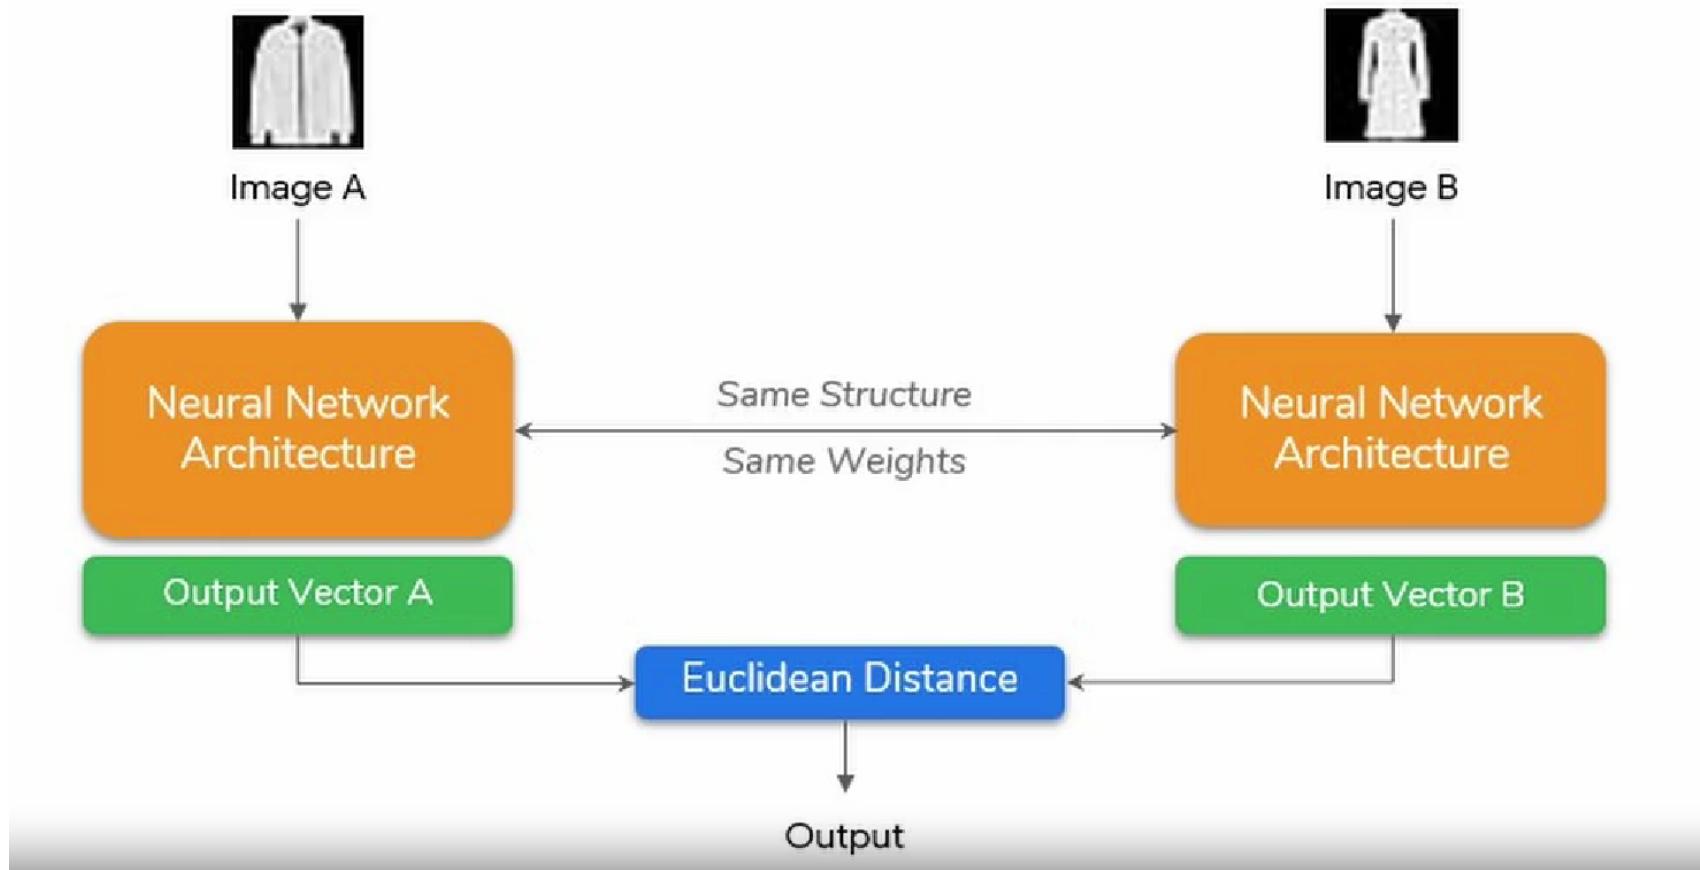
Based on the definition of the loss, there are three categories of triplets:

- **easy triplets**: triplets which have a loss of 0, because $d(a, p) + \text{margin} < d(a, n)$
- **hard triplets**: triplets where the negative is closer to the anchor than the positive, i.e. $d(a, n) < d(a, p)$
- **semi-hard triplets**: triplets where the negative is not closer to the anchor than the positive, but which still have positive loss: $d(a, p) < d(a, n) < d(a, p) + \text{margin}$

Each of these definitions depend on where the negative is, relatively to the anchor and positive. We can therefore extend these three categories to the negatives: **hard negatives**, **semi-hard negatives** or **easy negatives**.



A Siamese network's architecture



**Semantic segmentation,
Object localization,
Object detection**

Semantic Segmentation



GRASS, CAT,
TREE, SKY

No objects, just pixels

Classification + Localization



CAT

Single Object

Object Detection



DOG, DOG, CAT

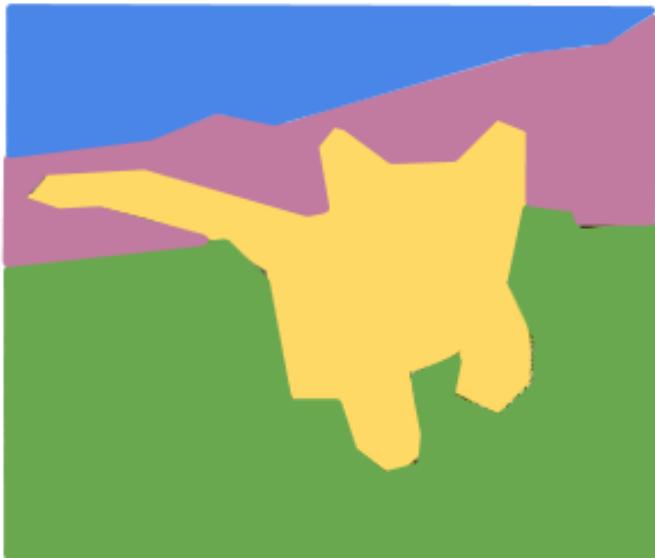
Multiple Object

Instance Segmentation



DOG, DOG, CAT

This image is CC0 public domain

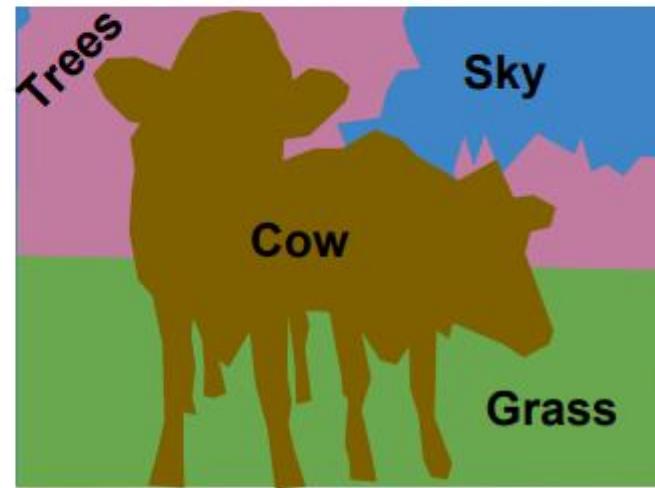
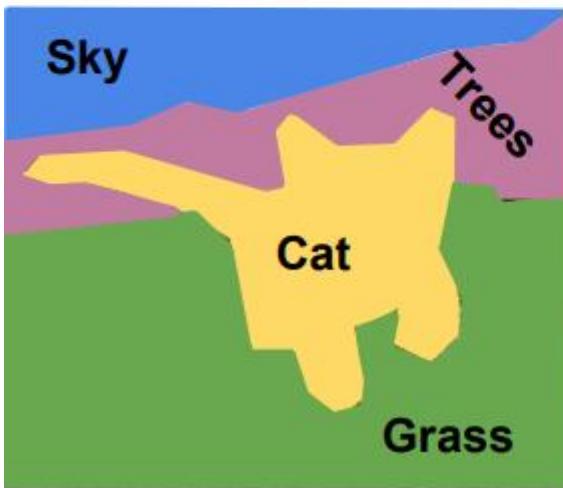


**GRASS, CAT,
TREE, SKY**

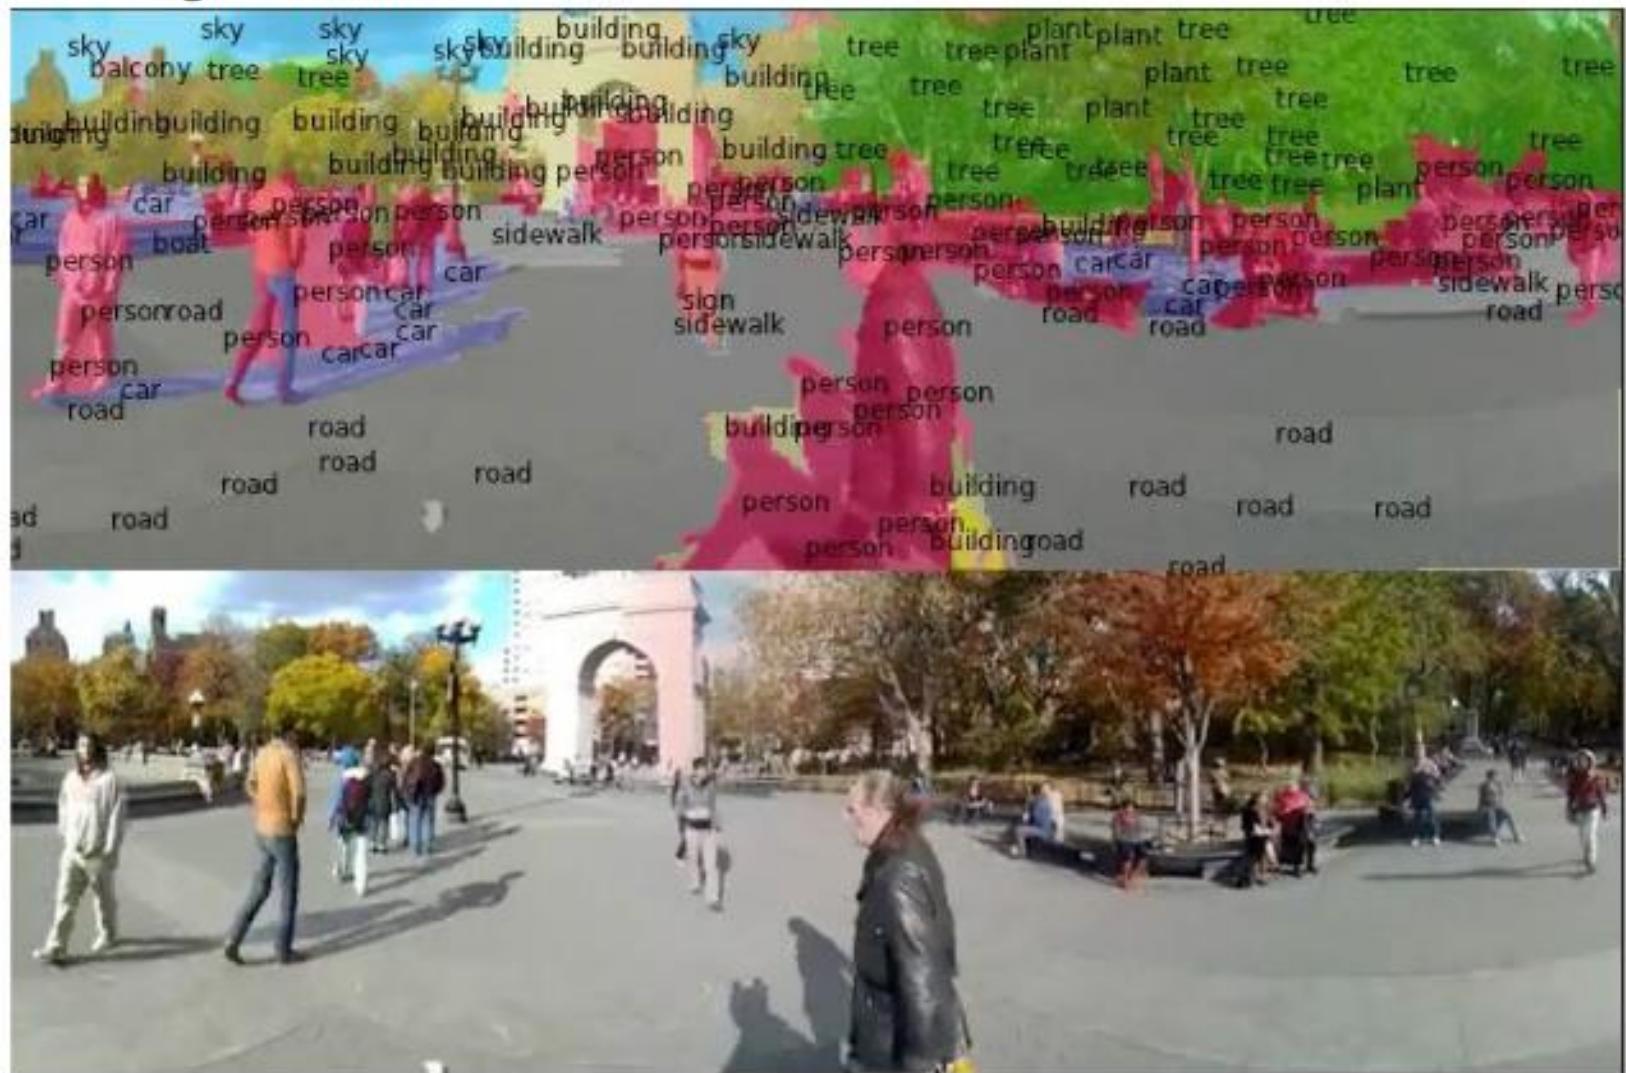
Semantic Segmentation



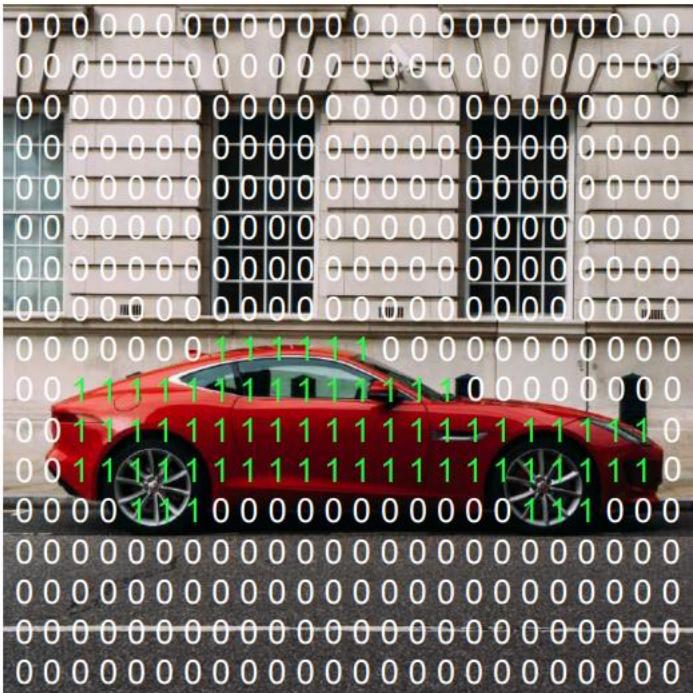
[This image is CC0 public domain](#)



Segmentation

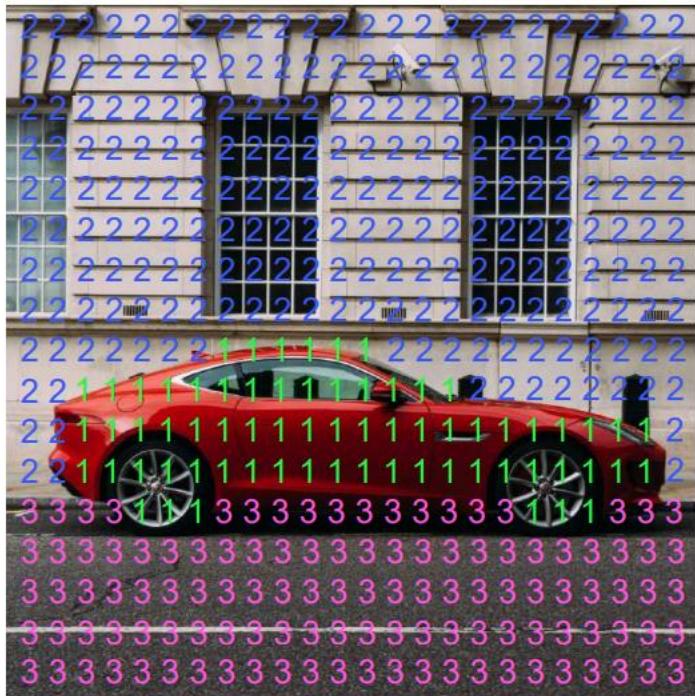


Per-pixel class labels



- 1. Car
- 0. Not Car

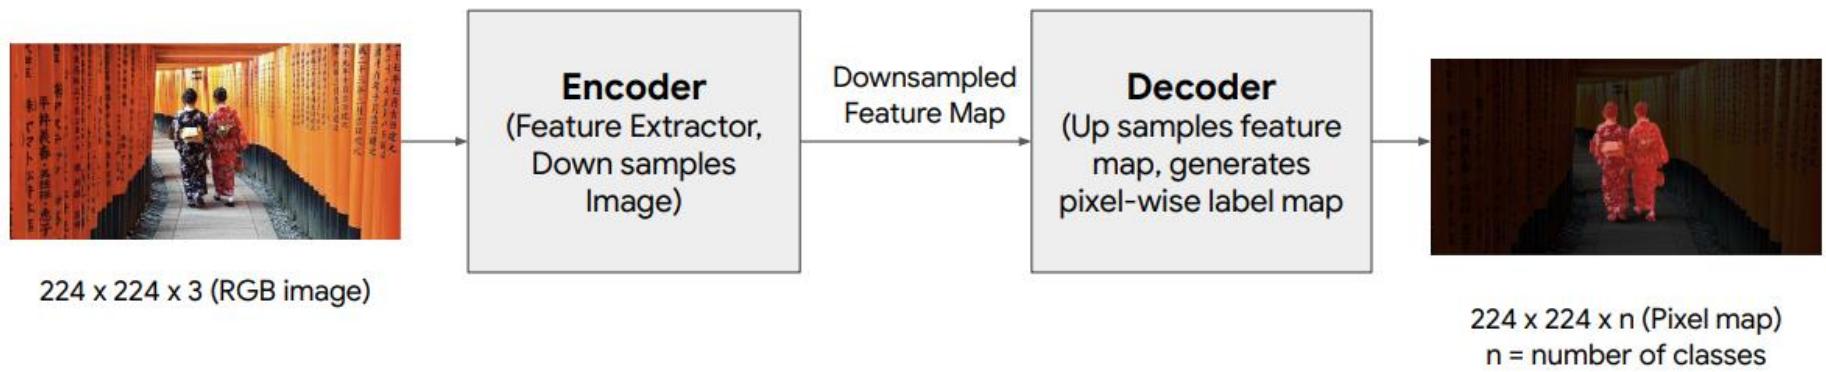
Per-pixel class labels



1. Car
 2. Building
 3. Road

Segmentation Map

Image Segmentation Basic Architecture



Semantic Segmentation Idea: Fully Convolutional

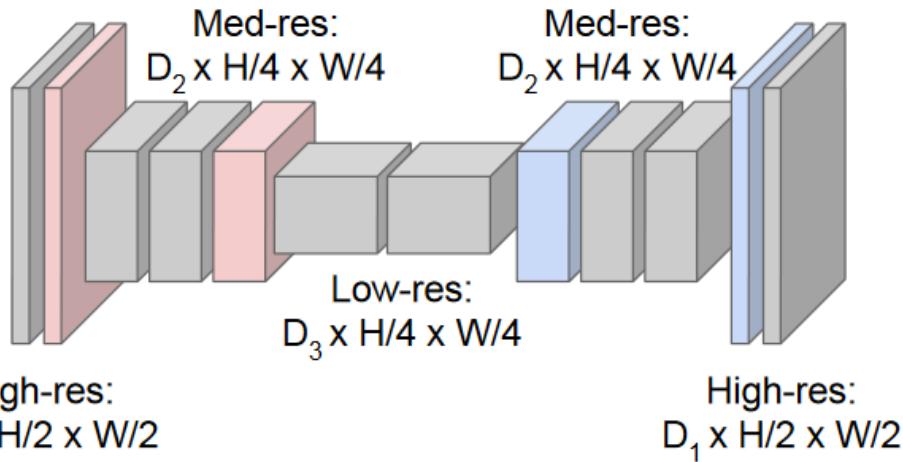
Downsampling:
Pooling, strided convolution



Input:
 $3 \times H \times W$

High-res:
 $D_1 \times H/2 \times W/2$

Design network as a bunch of convolutional layers, with
downsampling and **upsampling** inside the network!



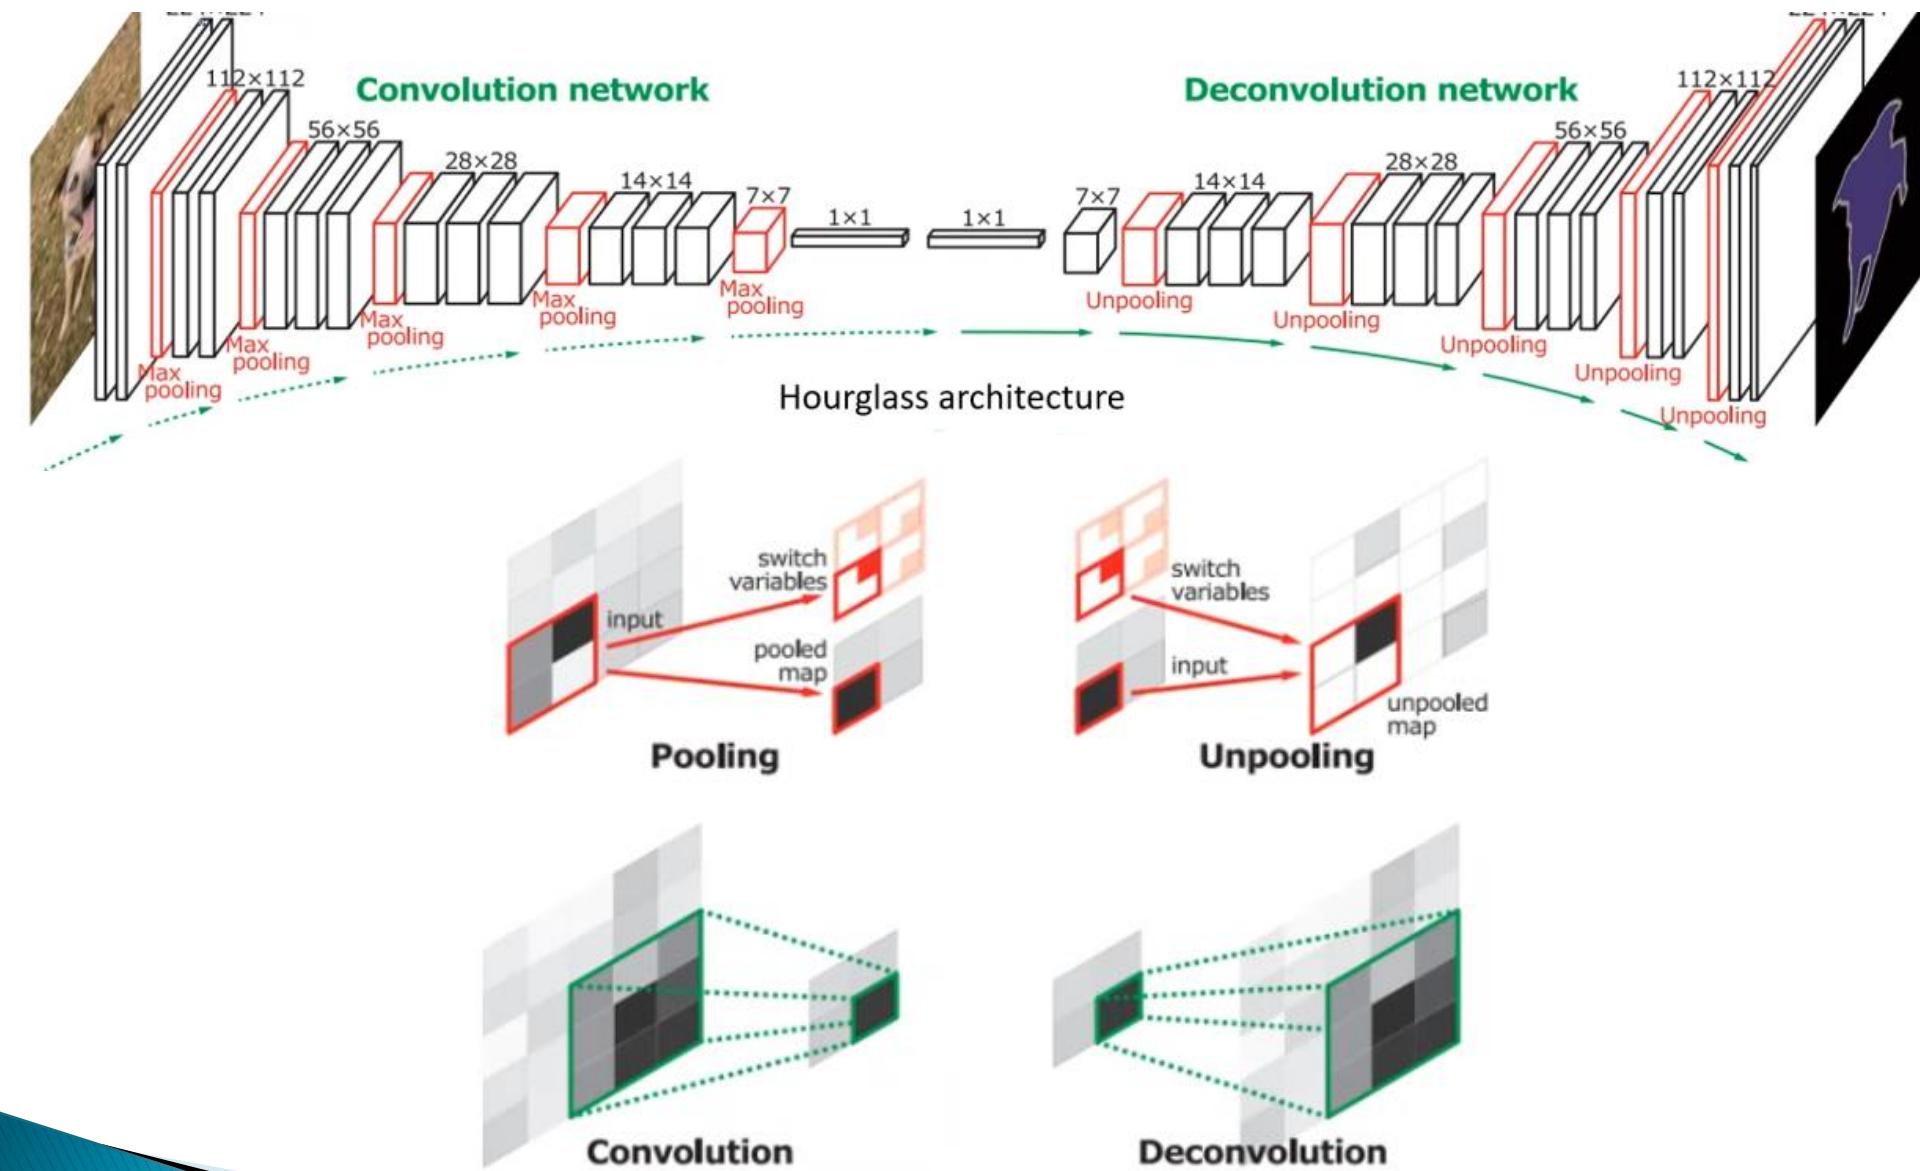
Upsampling:
Unpooling or strided transpose convolution



Predictions:
 $H \times W$

Long, Shelhamer, and Darrell, "Fully Convolutional Networks for Semantic Segmentation", CVPR 2015
Noh et al, "Learning Deconvolution Network for Semantic Segmentation", ICCV 2015

SegNet



In-Network upsampling: “Unpooling”

Nearest Neighbor

1	2
1	2
3	4



1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Input: 2 x 2

Output: 4 x 4

“Bed of Nails”

1	2
3	4



1	0	2	0
0	0	0	0
3	0	4	0
0	0	0	0

Input: 2 x 2

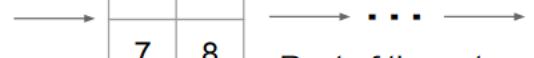
Output: 4 x 4

In-Network upsampling: “Max Unpooling”

Max Pooling

Remember which element was max!

1	2	6	3
3	5	2	1
1	2	2	1
7	3	4	8



Input: 4 x 4

Output: 2 x 2

Max Unpooling

Use positions from
pooling layer

1	2
3	4



0	0	2	0
0	1	0	0
0	0	0	0
3	0	0	4

Input: 2 x 2

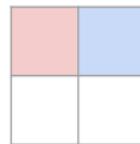
Output: 4 x 4

Learnable Upsampling: Transpose Convolution

Other names:

- Deconvolution (bad)
- Upconvolution
- Fractionally strided convolution
- Backward strided convolution

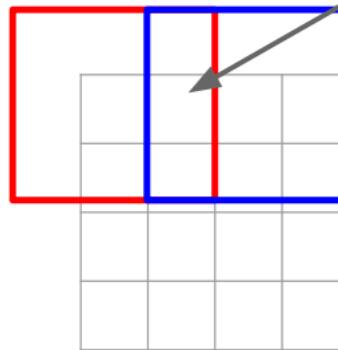
3 x 3 transpose convolution, stride 2 pad 1



Input: 2 x 2



Input gives weight for filter



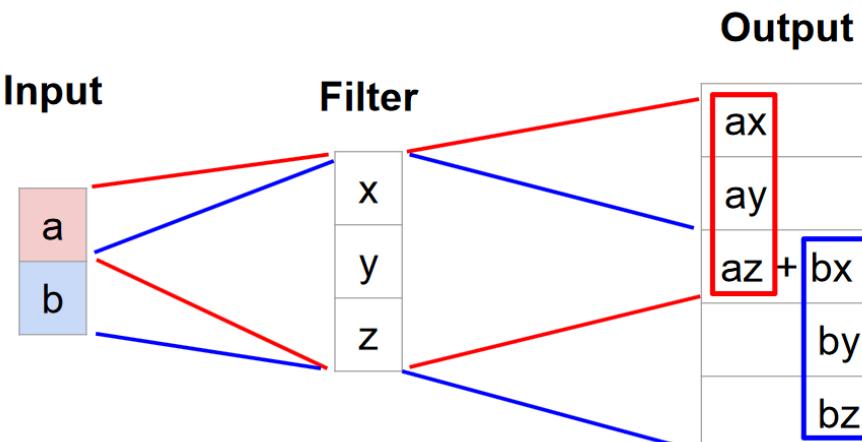
Output: 4 x 4

Sum where output overlaps

Filter moves 2 pixels in the output for every one pixel in the input

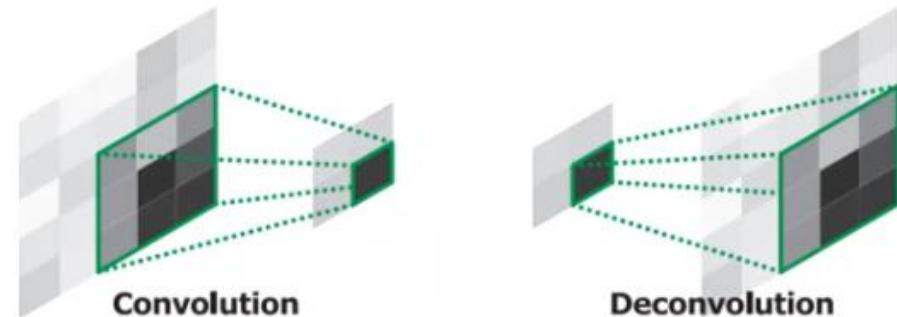
Stride gives ratio between movement in output and input

Transpose Convolution: 1D Example



Output

Output contains copies of the filter weighted by the input, summing at where at overlaps in the output



Filter

BATCHNORMULATION

Bidirectional

CategoryEncoding

CenterCrop

Concatenate

Conv1D

Conv1DTranspose

Conv2D

Conv2DTranspose

Conv3D

Conv3DTranspose

ConvLSTM1D

ConvLSTM2D

ConvLSTM3D

Cropping1D

Cropping2D

Cropping3D

Dense

DenseFeatures

DepthwiseConv1D

DepthwiseConv2D

Discretization

Dot

Dropout

Ellipsis

tf.keras.layers.Conv2DTranspose



[View source on GitHub](#)

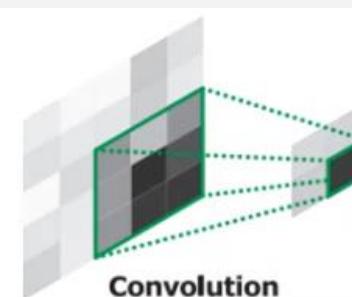
Transposed convolution layer (sometimes called Deconvolution).

Inherits From: [Conv2D](#), [Layer](#), [Module](#)

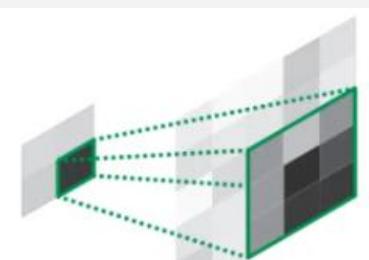
[View aliases](#)



```
tf.keras.layers.Conv2DTranspose(  
    filters,  
    kernel_size,  
    strides=(1, 1),  
    padding='valid',  
    output_padding=None,  
    data_format=None,  
    dilation_rate=(1, 1),  
    activation=None,  
    use_bias=True,  
    kernel_initializer='glorot_uniform',  
    bias_initializer='zeros',  
    kernel_regularizer=None,  
    bias_regularizer=None,  
    activity_regularizer=None,  
    kernel_constraint=None,  
    bias_constraint=None,  
    **kwargs
```



Convolution



Deconvolution

StackedRNNCells

StringLookup

Subtract

TextVectorization

ThresholdedReLU

TimeDistributed

UnitNormalization

UpSampling1D

UpSampling2D

UpSampling3D

Wrapper

ZeroPadding1D

ZeroPadding2D

ZeroPadding3D

add

average

concatenate

deserialize

dot

maximum

minimum

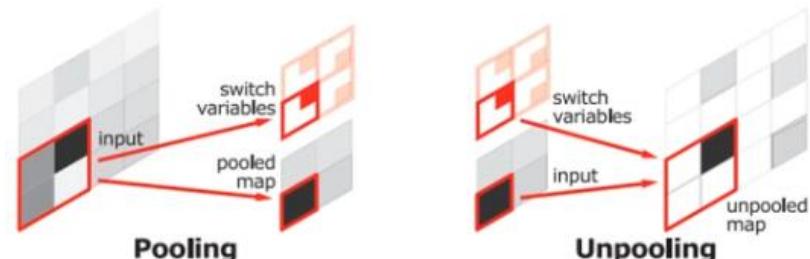
multiply

serialize

tf.keras.layers.UpSampling2D

[View source on GitHub](#)

Upsampling layer for 2D inputs.

Inherits From: [Layer](#), [Module](#)[+ View aliases](#)

```
tf.keras.layers.UpSampling2D(  
    size=(2, 2), data_format=None, interpolation='nearest', **kwargs  
)
```



Used in the notebooks

Args

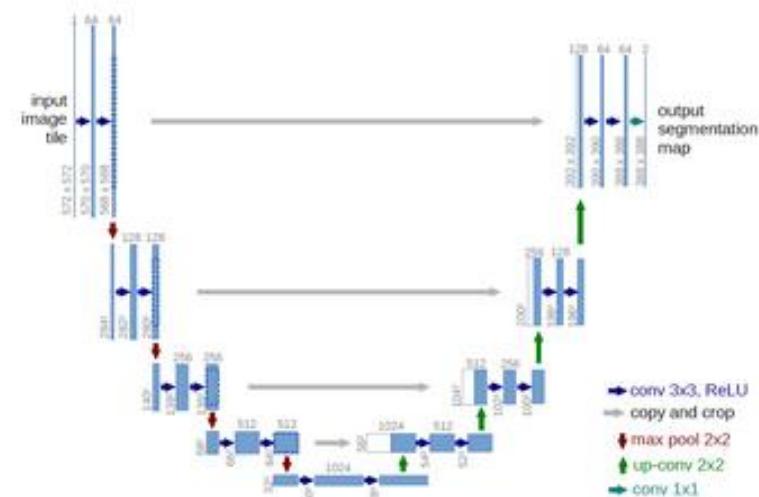
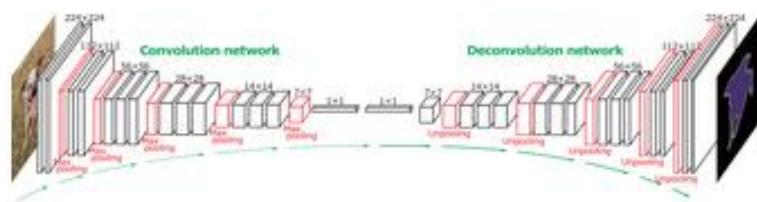
size	Int, or tuple of 2 integers. The upsampling factors for rows and columns.
-------------	---

data_format	A string, one of <code>channels_last</code> (default) or <code>channels_first</code> . The ordering of the dimensions in the inputs. <code>channels_last</code> corresponds to inputs with shape <code>(batch_size, height, width, channels)</code> while <code>channels_first</code> corresponds to inputs with shape <code>(batch_size, channels, height, width)</code> . It defaults to the <code>image_data_format</code> value found in your Keras config file at <code>~/.keras/keras.json</code> . If you never set it, then it will be "channels_last".
--------------------	---

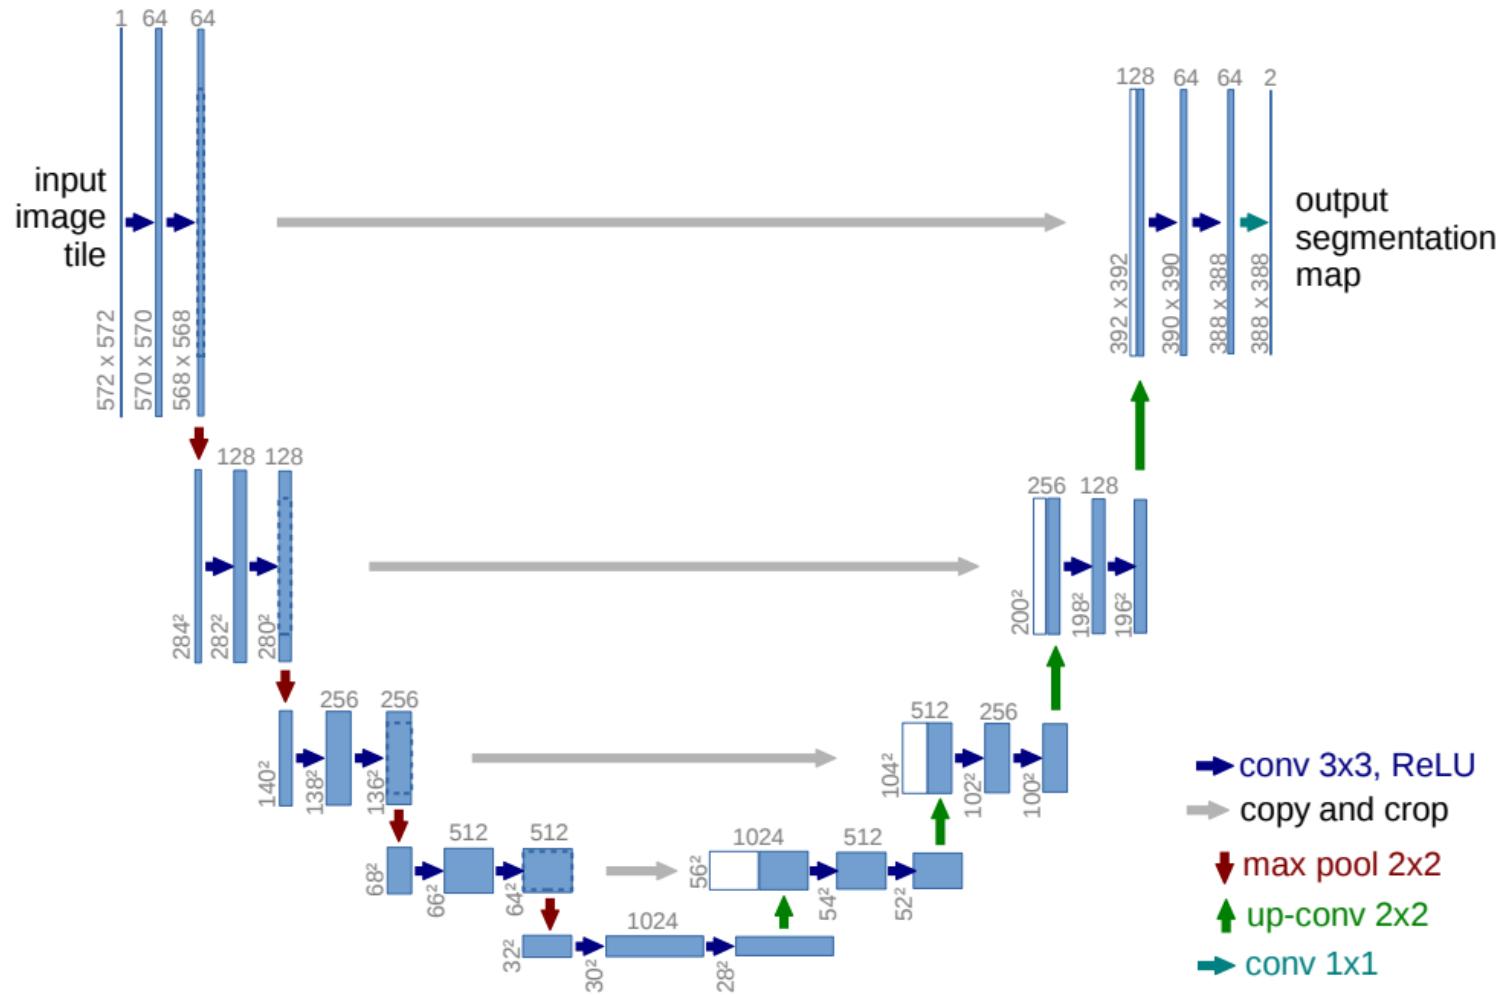
interpolation	A string, one of "area", "bicubic", "bilinear", "gaussian", "lanczos3", "lanczos5", "mitchellcubic", "nearest".
----------------------	---

UNet

SegNet



U-net



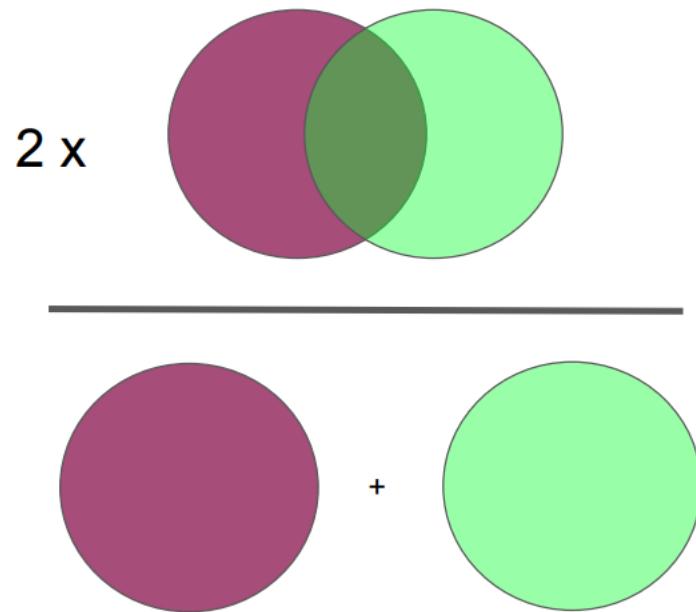
Metrics

Table 1. The three similarity coefficients

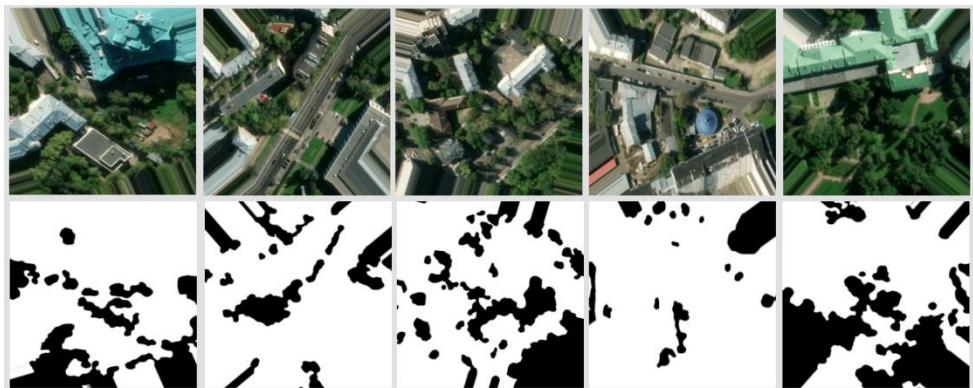
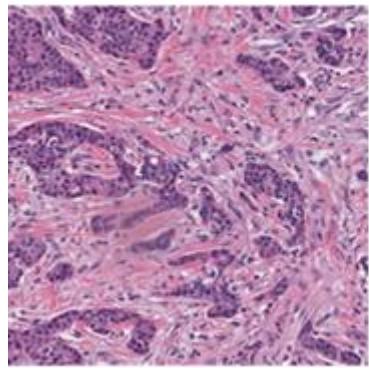
Similarity Coefficient (X,Y)	Actual Formula
Dice Coefficient	$\frac{2 X \cap Y }{ X + Y }$
Cosine Coefficient	$\frac{ X \cap Y }{ X ^{1/2} \cdot Y ^{1/2}}$
Jaccard Coefficient	$\frac{ X \cap Y }{ X + Y - X \cap Y }$

Dice Score

$$\text{Dice Score} = 2 \times \frac{\text{Area of Overlap}}{\text{Combined Area}}$$



Practical applications



(a) original image



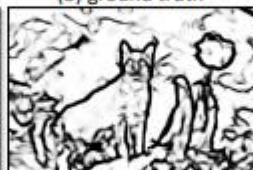
(b) ground truth



(c) HED: output



(d) HED: side output 2



(e) HED: side output 3



(f) HED: side output 4

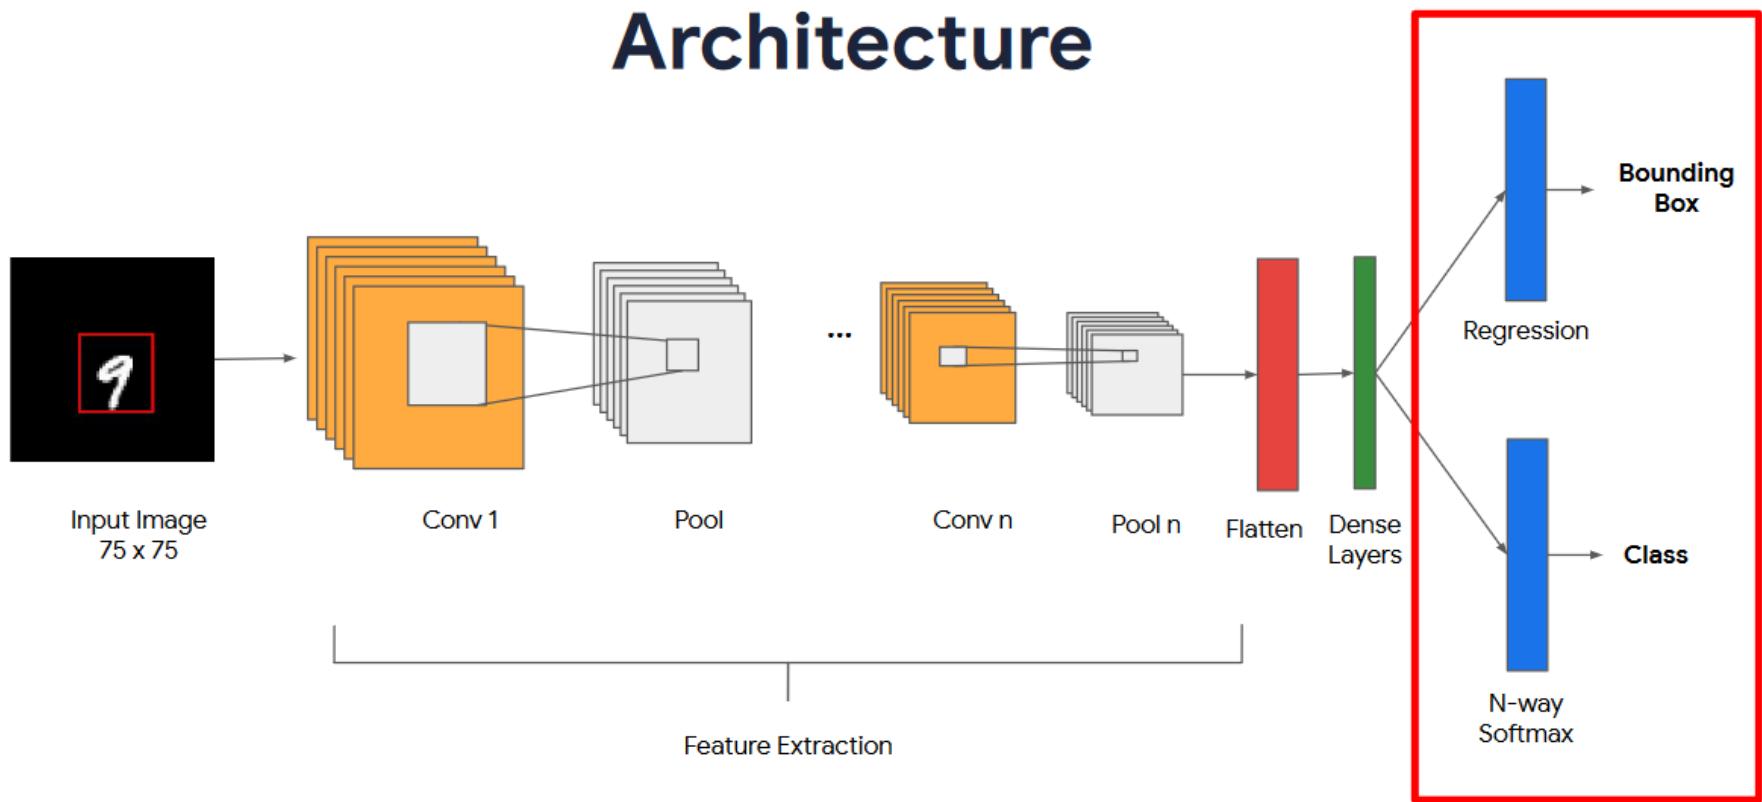


Classification + Localization



CAT

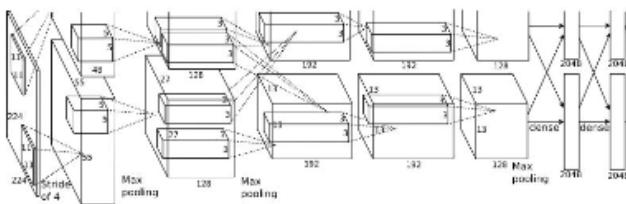
Convolutional Neural Networks Architecture



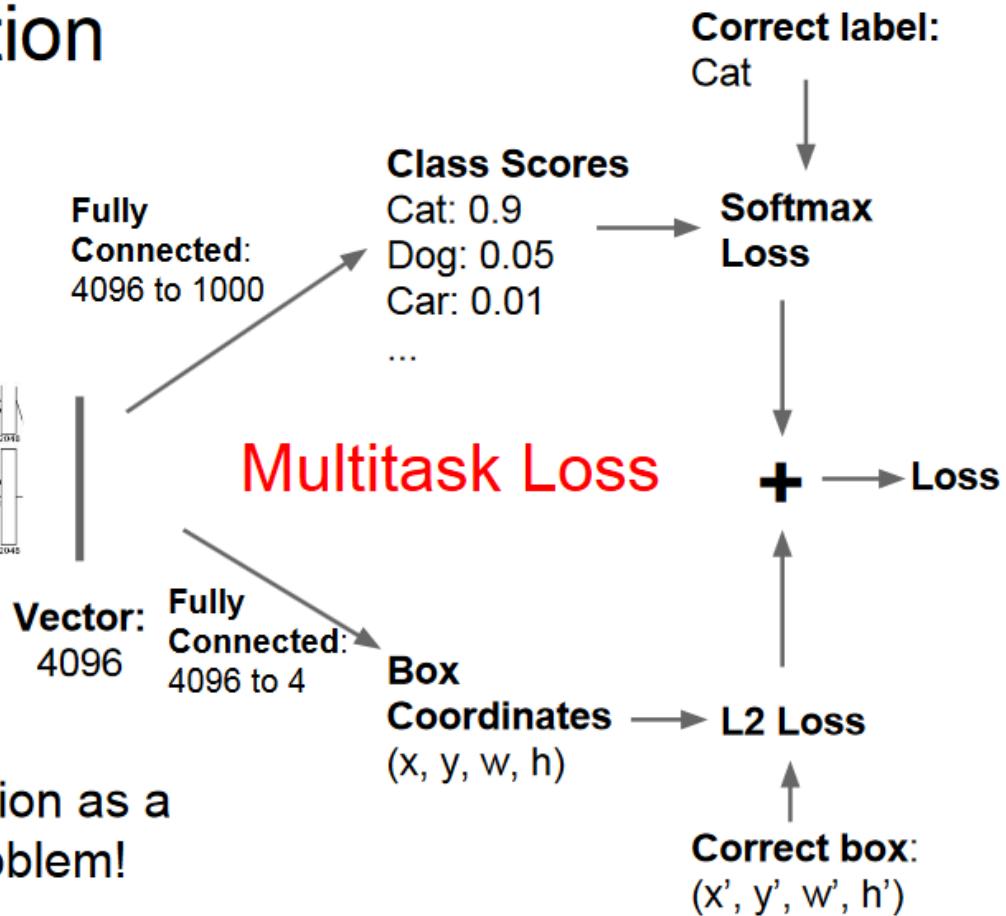
Classification + Localization



This image is CC0 public domain



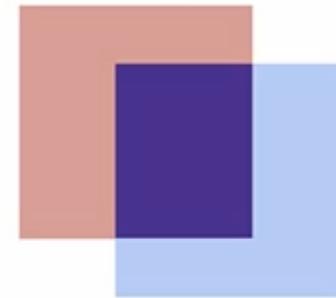
Treat localization as a
regression problem!



Intersection Over Union

IOU =

Intersection

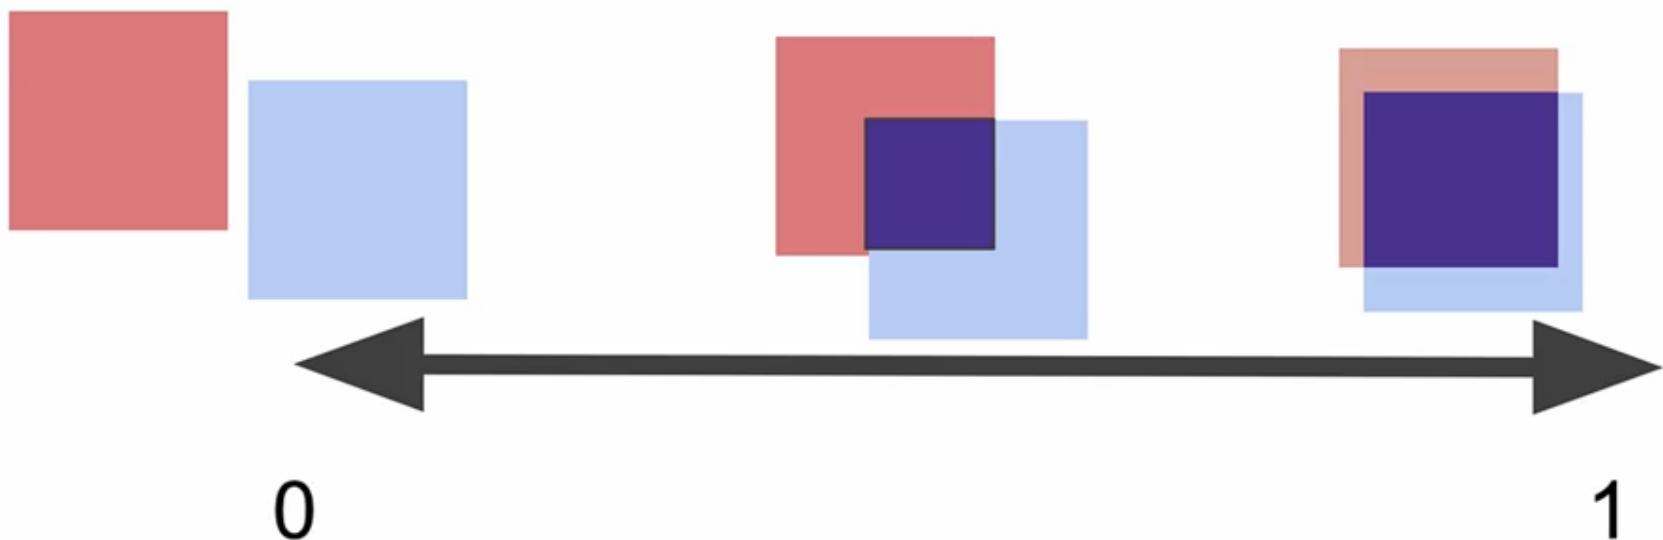


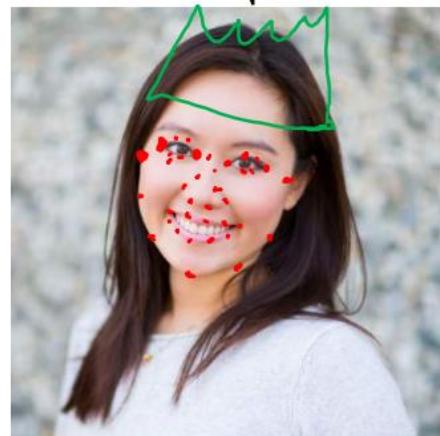
Union



Intersection Over Union

$$\text{IOU} = \frac{\text{Intersection}}{\text{Union}}$$



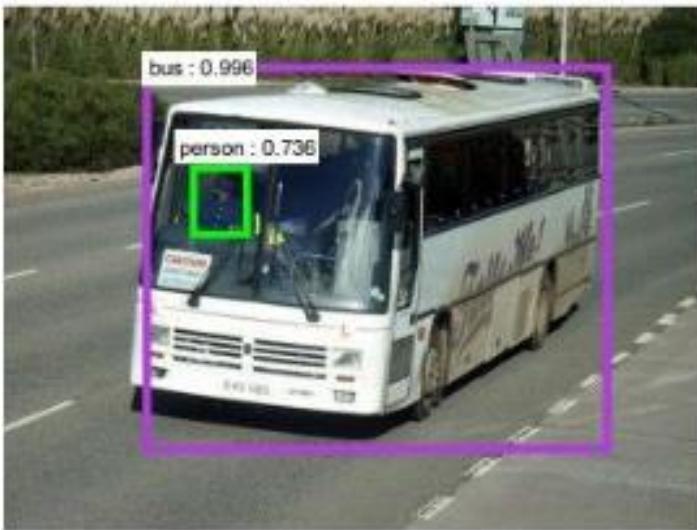
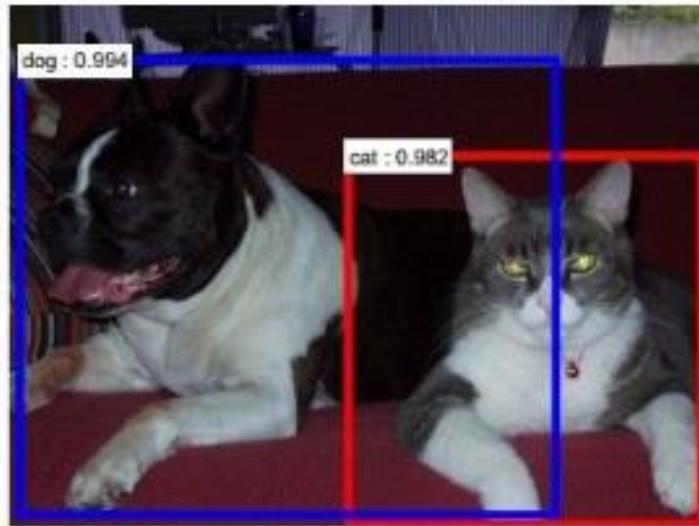
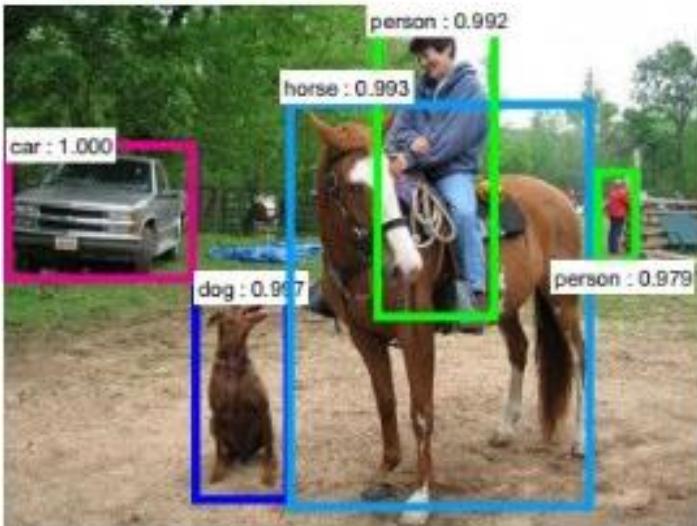




DOG, DOG, CAT

Object Detection

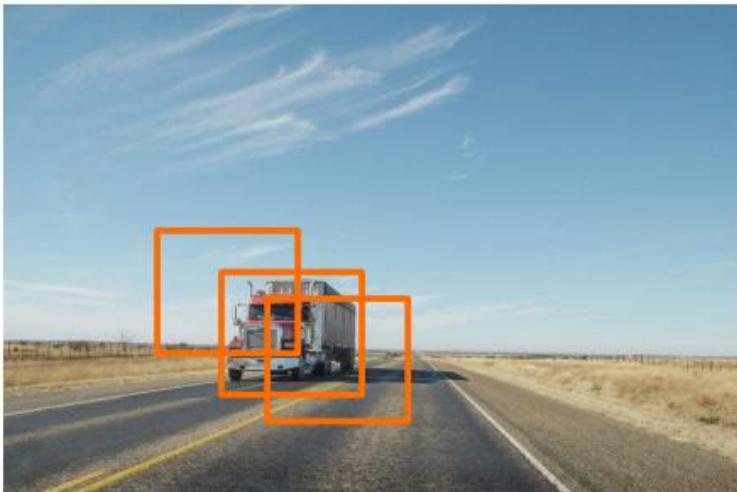
Detection



Two stages to object detection

- ▶ Region proposal
- ▶ Object detection and classification

Non-maximum suppression (NMS)



Before NMS



After NMS

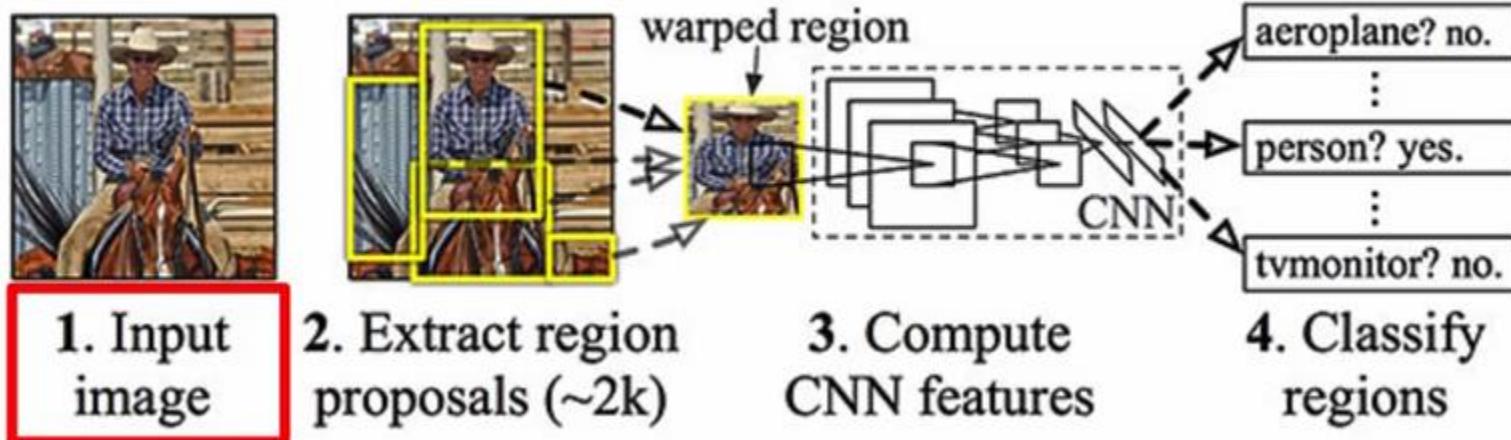
Object detection methods

One-shot
(YOLO, SSD)

Two-shot
(R-CNN, Fast RCNN,
Faster RCNN)

R-CNN

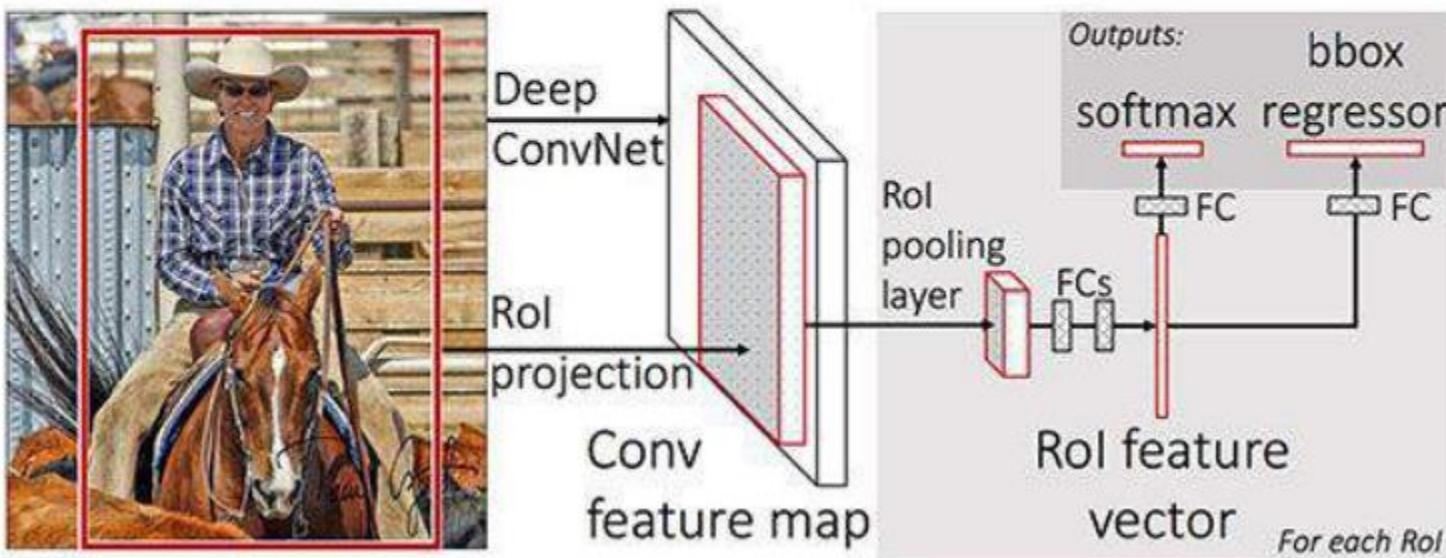
R-CNN: *Regions with CNN features*

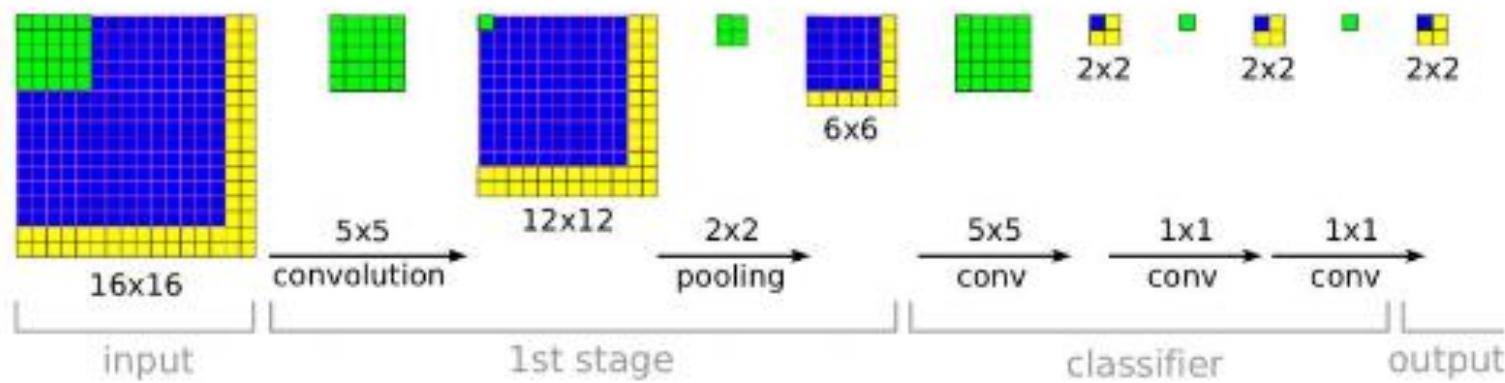
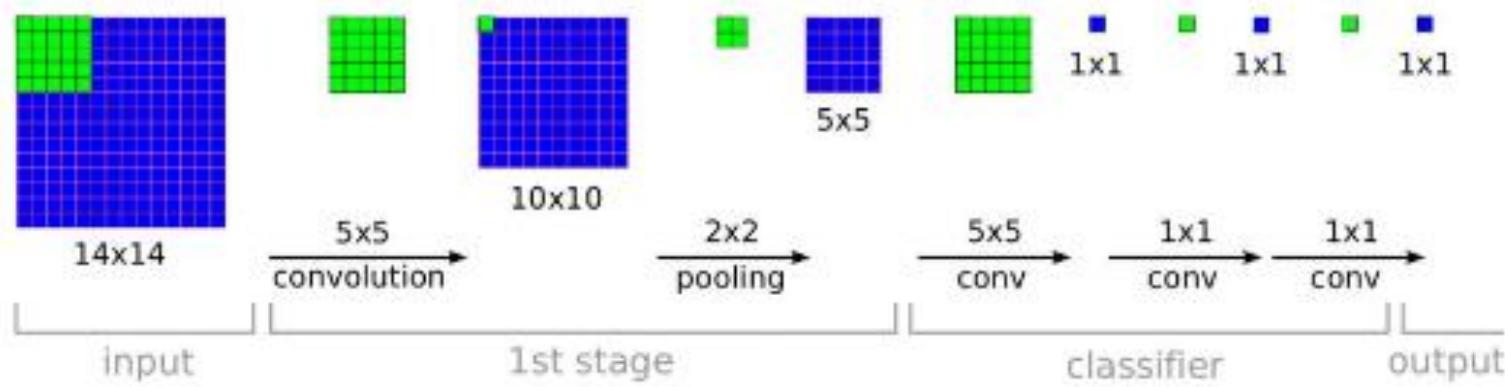


<https://arxiv.org/abs/1311.2524>

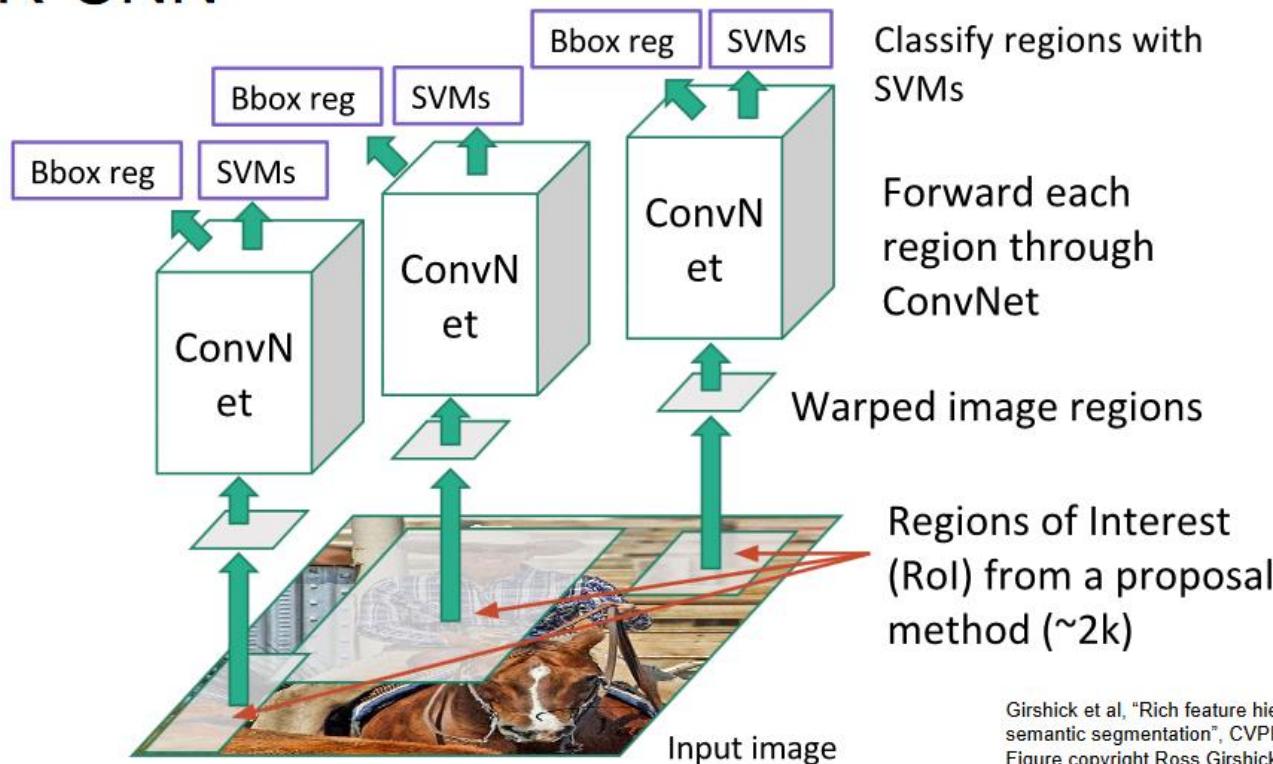
<https://www.koen.me/research/pub/uijlings-ijcv2013-draft.pdf>

Fast R-CNN





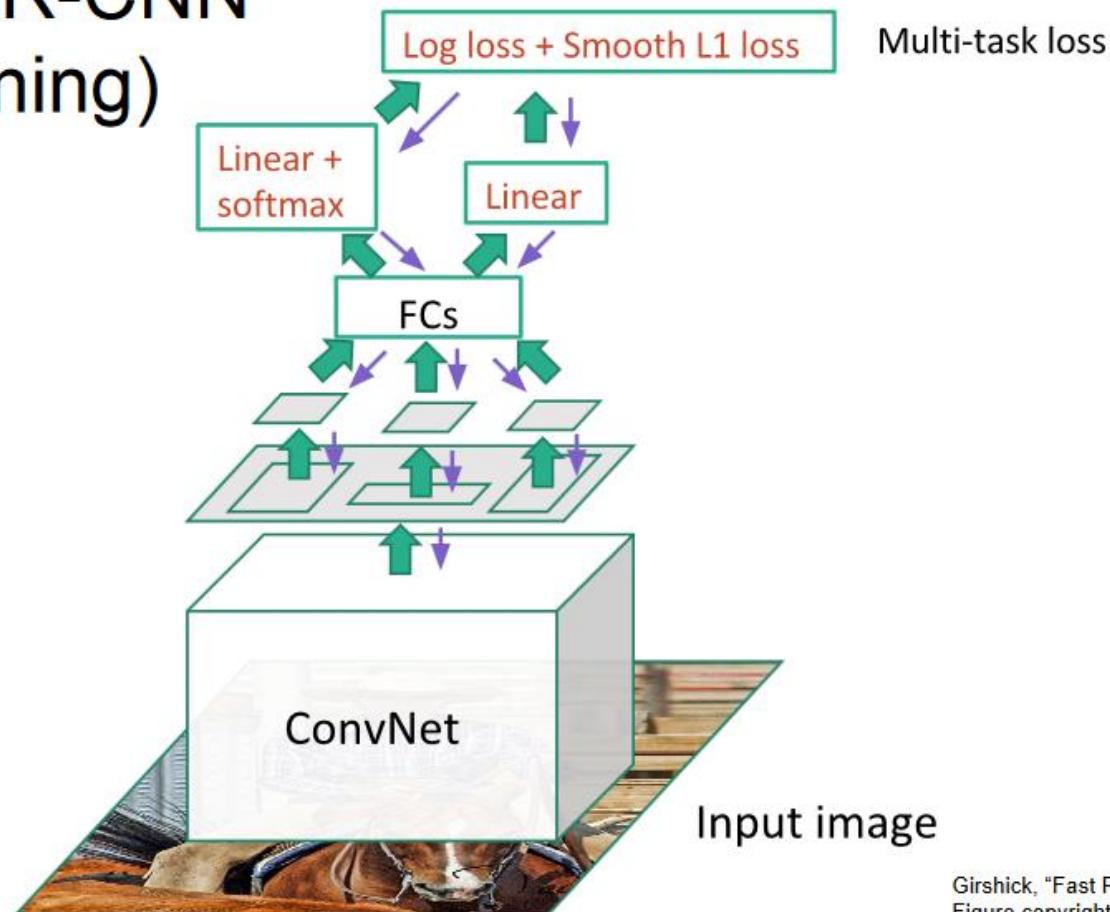
R-CNN



Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.

Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

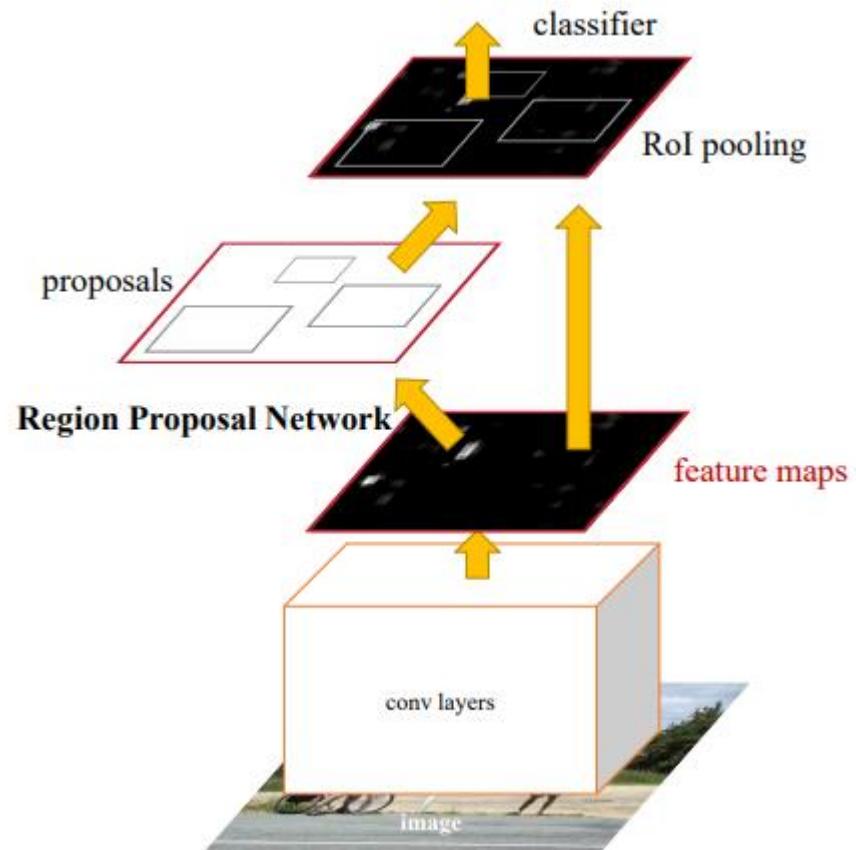
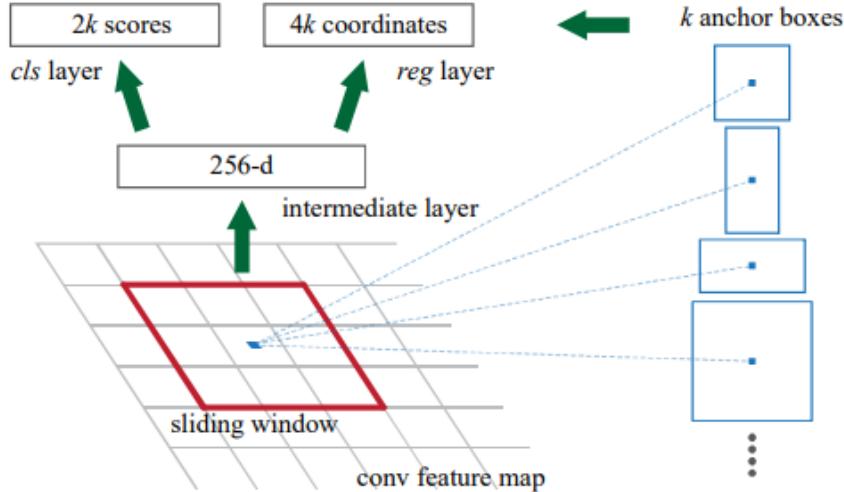
Fast R-CNN (Training)



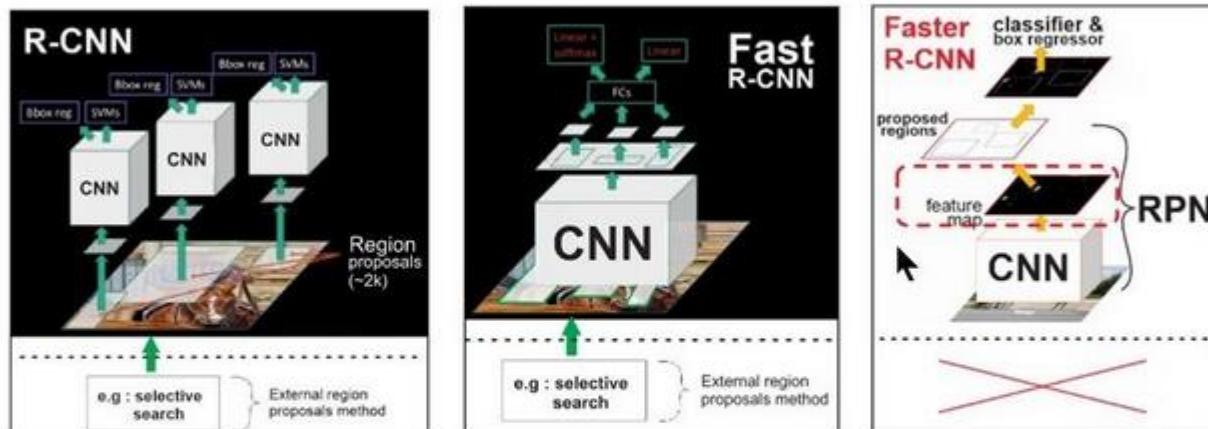
Girshick, "Fast R-CNN", ICCV 2015.

Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

Faster R-CNN



Two shot: performance

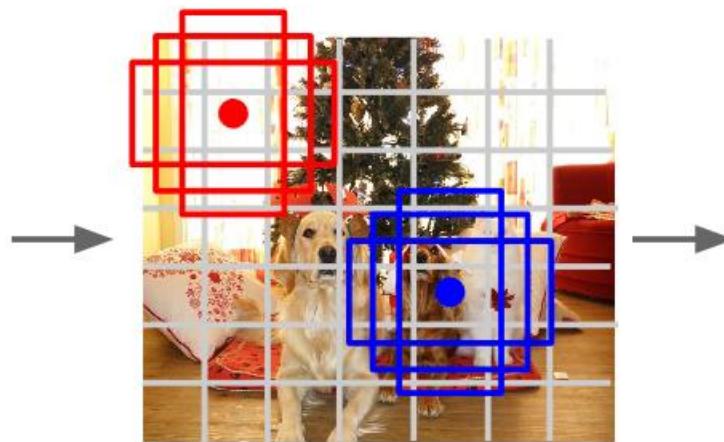


	R-CNN	Fast R-CNN	Faster R-CNN
Test time per image	50 seconds	2 seconds	0.2 seconds
Speed-up	1x	25x	250x
mAP (VOC 2007)	66.0%	66.9%	66.9%

Detection without Proposals: YOLO / SSD



Input image
 $3 \times H \times W$



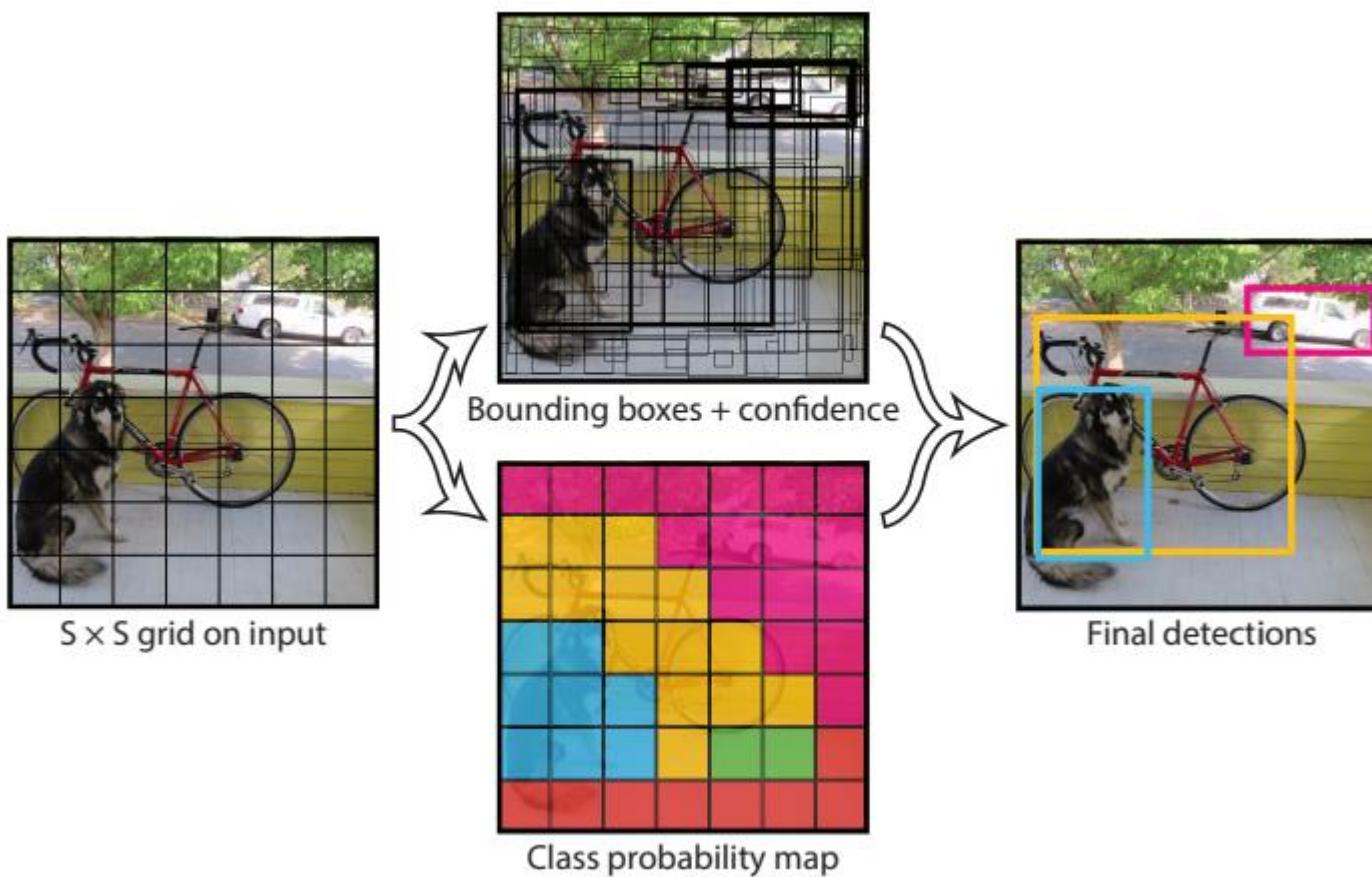
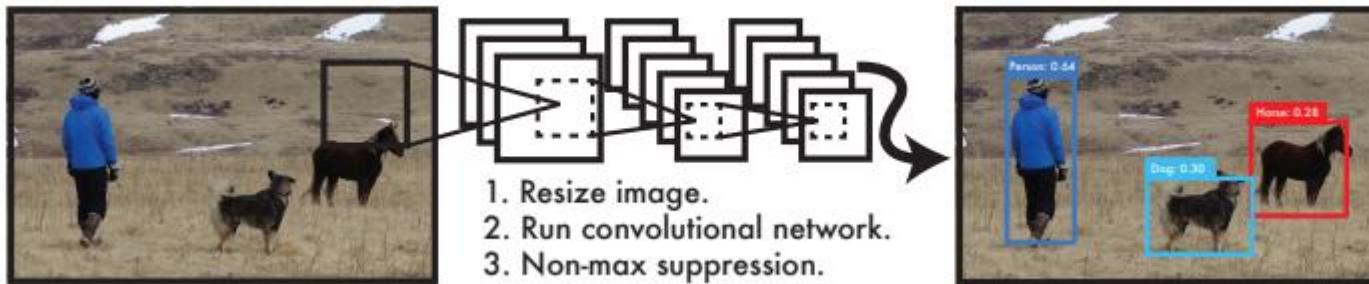
Divide image into grid
 7×7

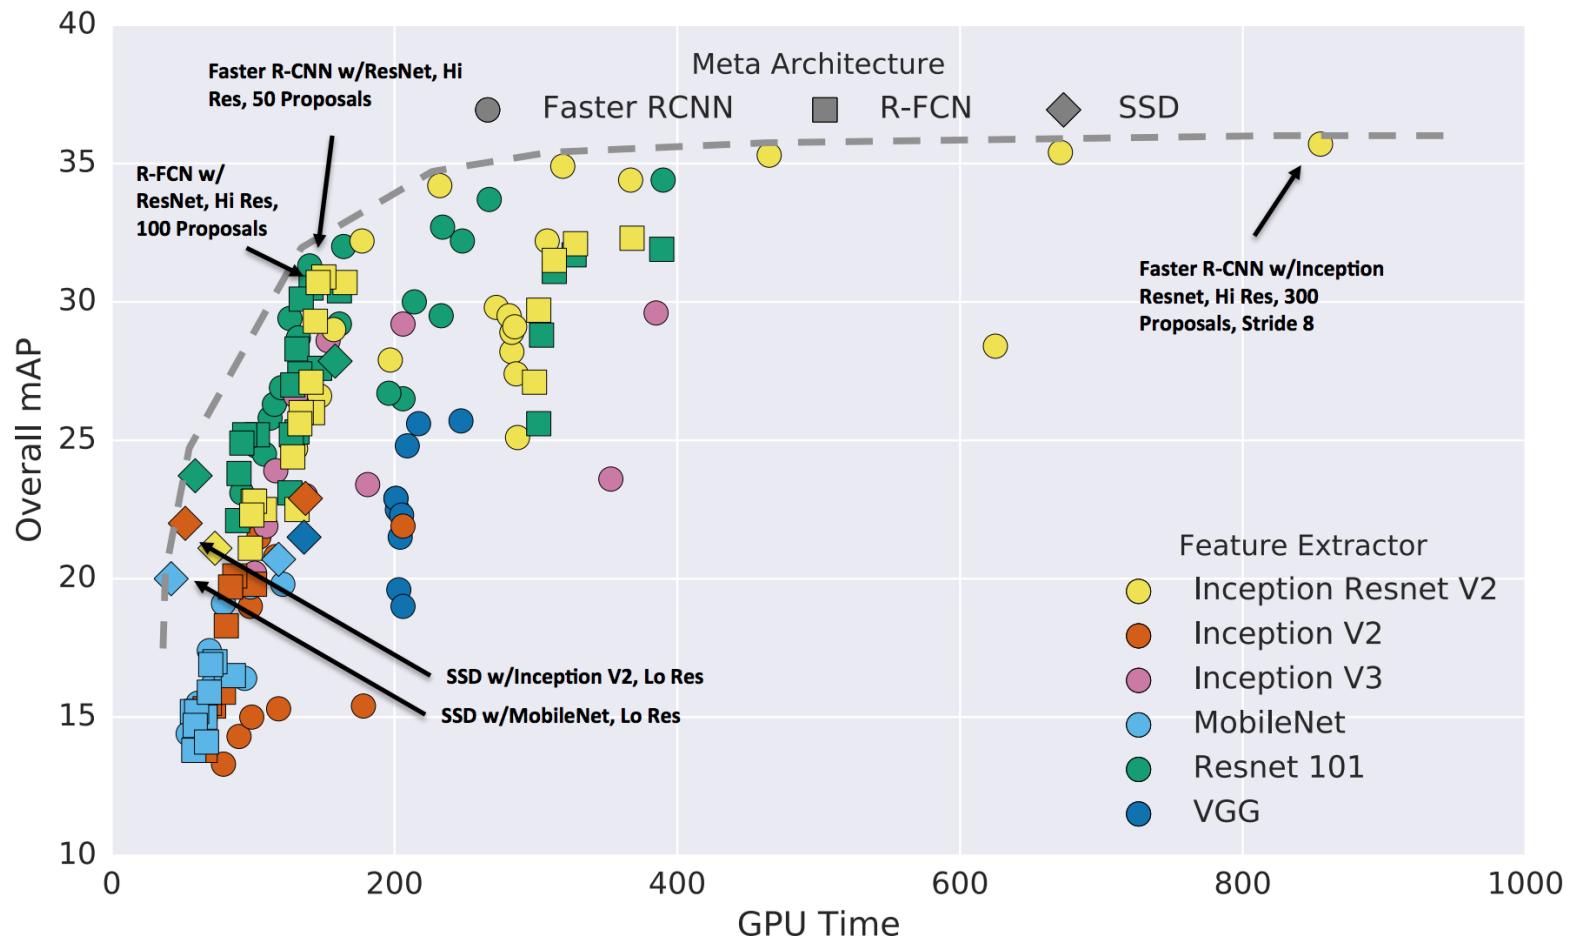
Image a set of **base boxes**
centered at each grid cell
Here $B = 3$

- Within each grid cell:
- Regress from each of the B base boxes to a final box with 5 numbers:
(dx , dy , dh , dw , confidence)
 - Predict scores for each of C classes (including background as a class)

Output:
 $7 \times 7 \times (5 * B + C)$

Redmon et al, "You Only Look Once: Unified, Real-Time Object Detection", CVPR 2016
Liu et al, "SSD: Single-Shot MultiBox Detector", ECCV 2016





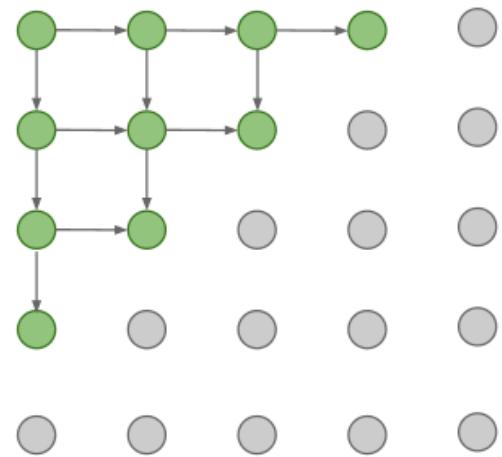
Generative models

PixelRNN

[van der Oord et al. 2016]

Generate image pixels starting from corner

Dependency on previous pixels modeled using an RNN (LSTM)



PixelCNN

[van der Oord et al. 2016]

Still generate image pixels starting from corner

Dependency on previous pixels now modeled using a CNN over context region

Training is faster than PixelRNN
(can parallelize convolutions since context region values known from training images)

Generation must still proceed sequentially
=> still slow

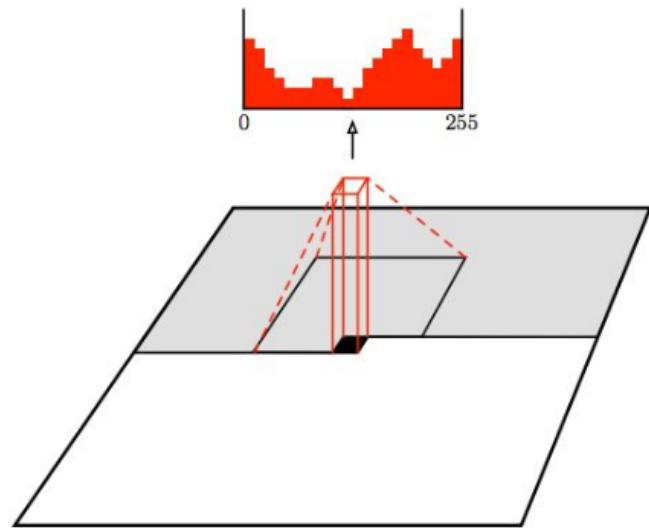


Figure copyright van der Oord et al., 2016. Reproduced with permission.

Autoencoders

Train such that features can be used to reconstruct original data

Reconstructed input data

Features

Input data

L2 Loss function:

$$\|x - \hat{x}\|^2$$

\hat{x}

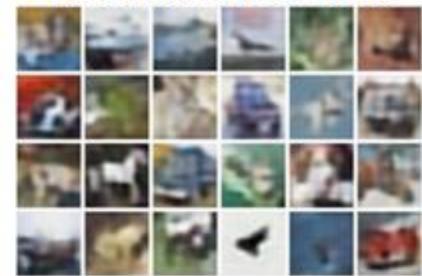
Decoder

z

Encoder

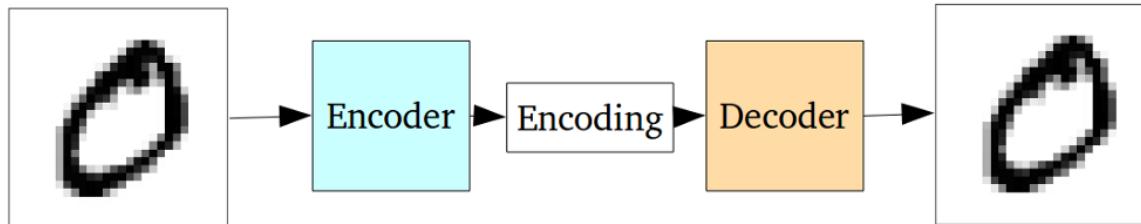
x

Reconstructed data

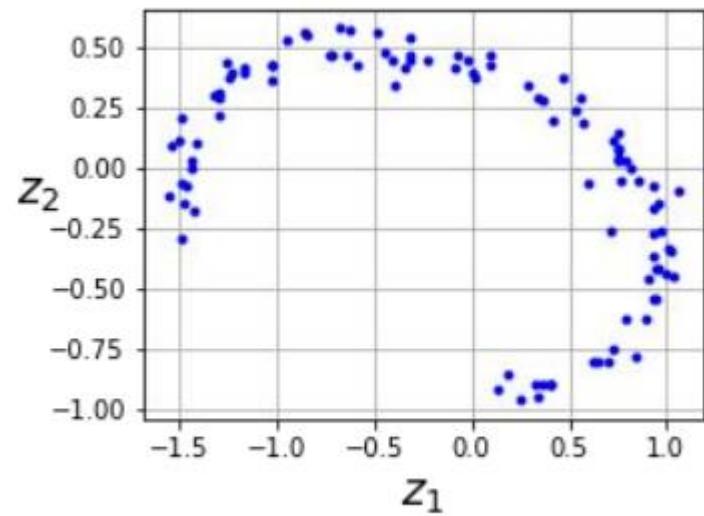
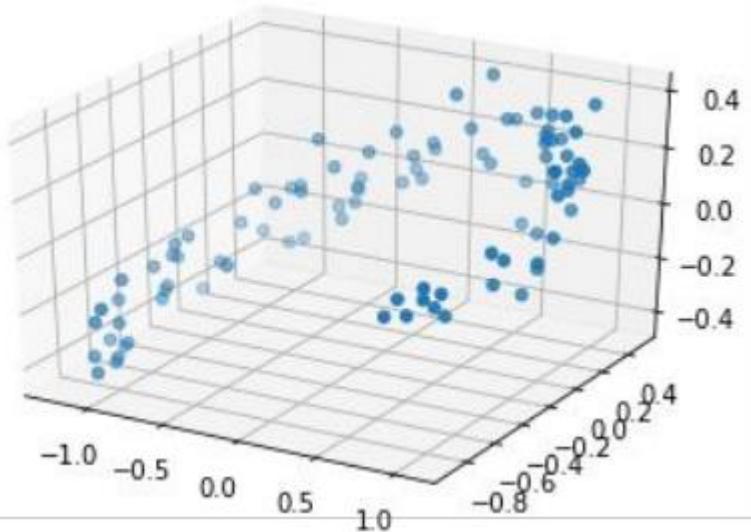


Encoder: 4-layer conv
Decoder: 4-layer upconv

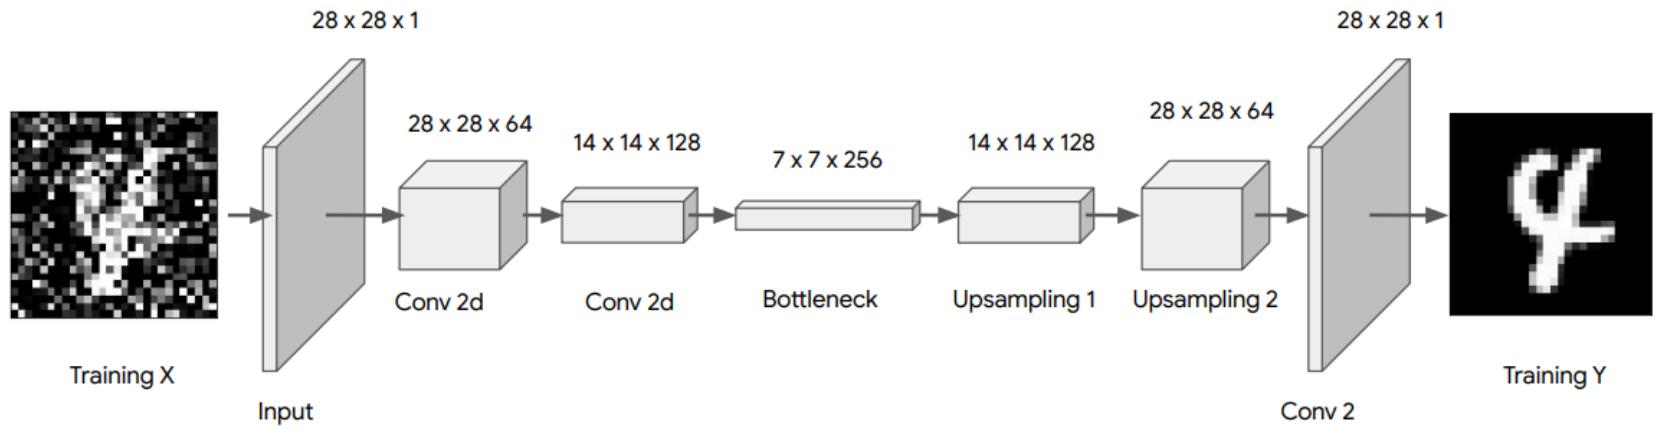
Input data



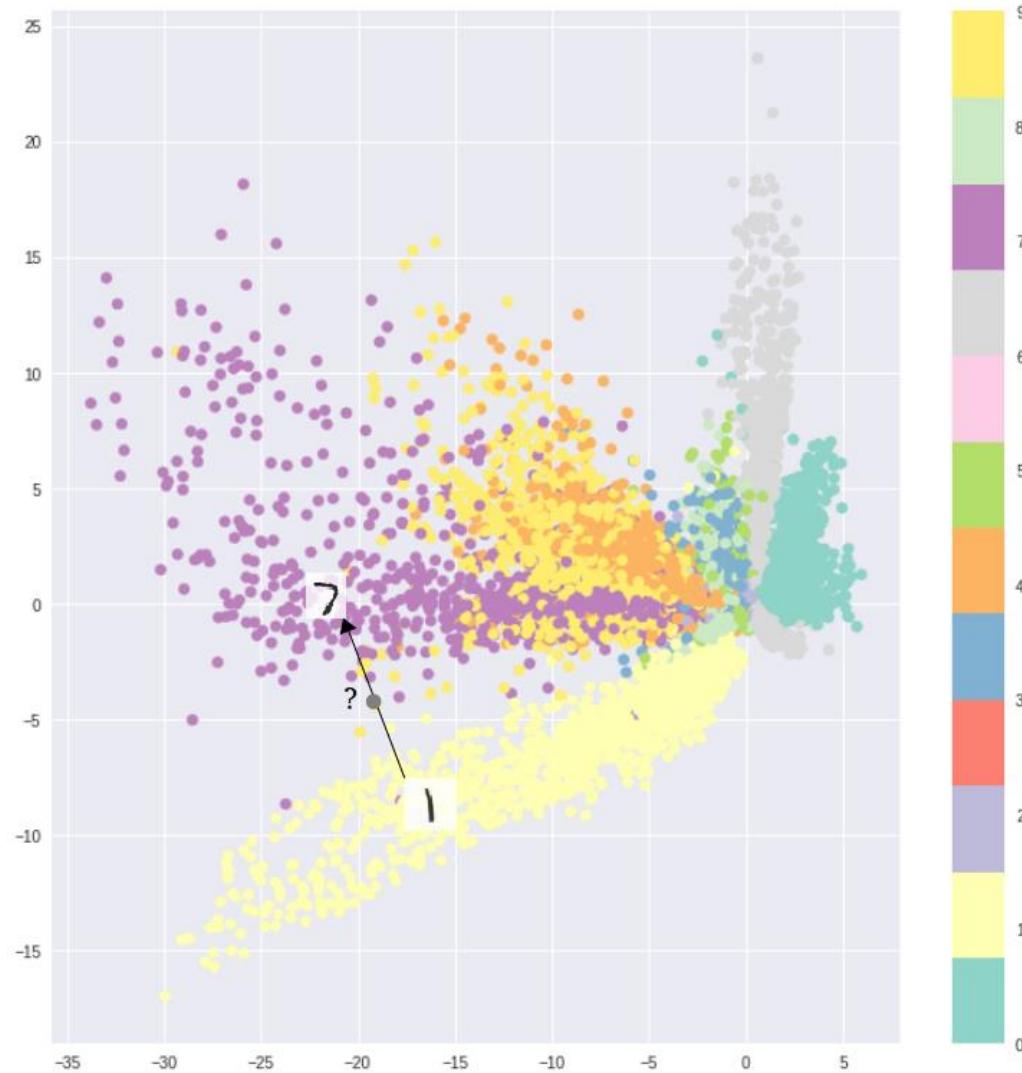
Dimensionality reduction



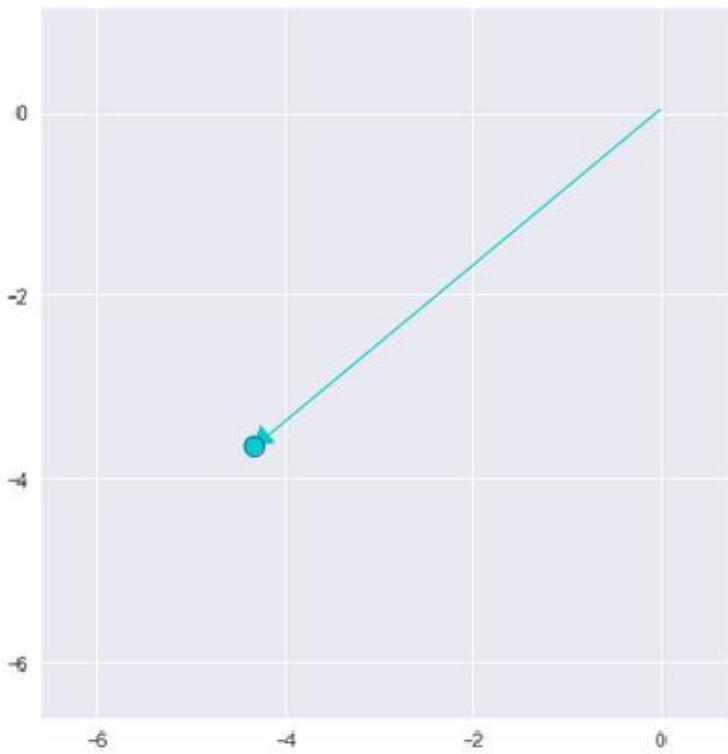
Noise reduction



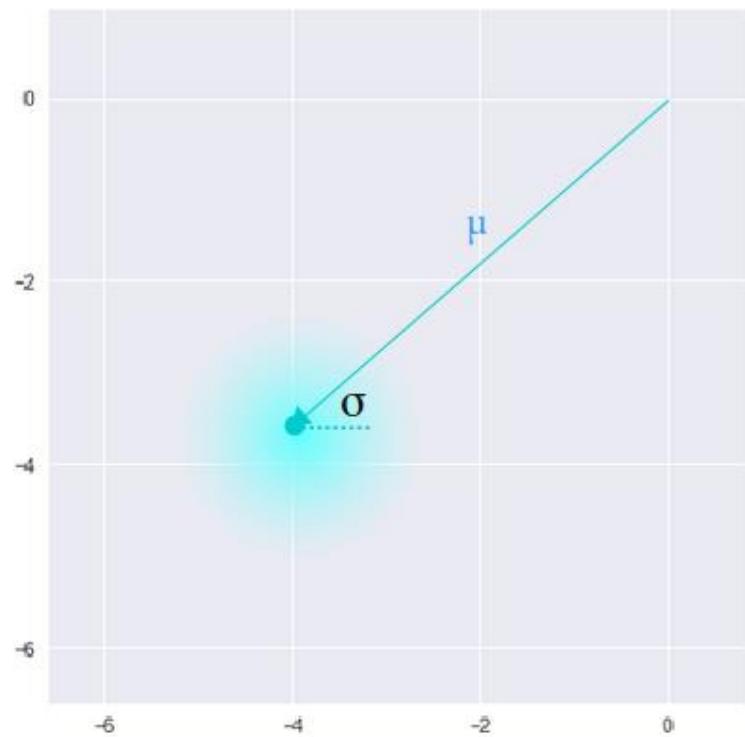
Autoencoder as generative model



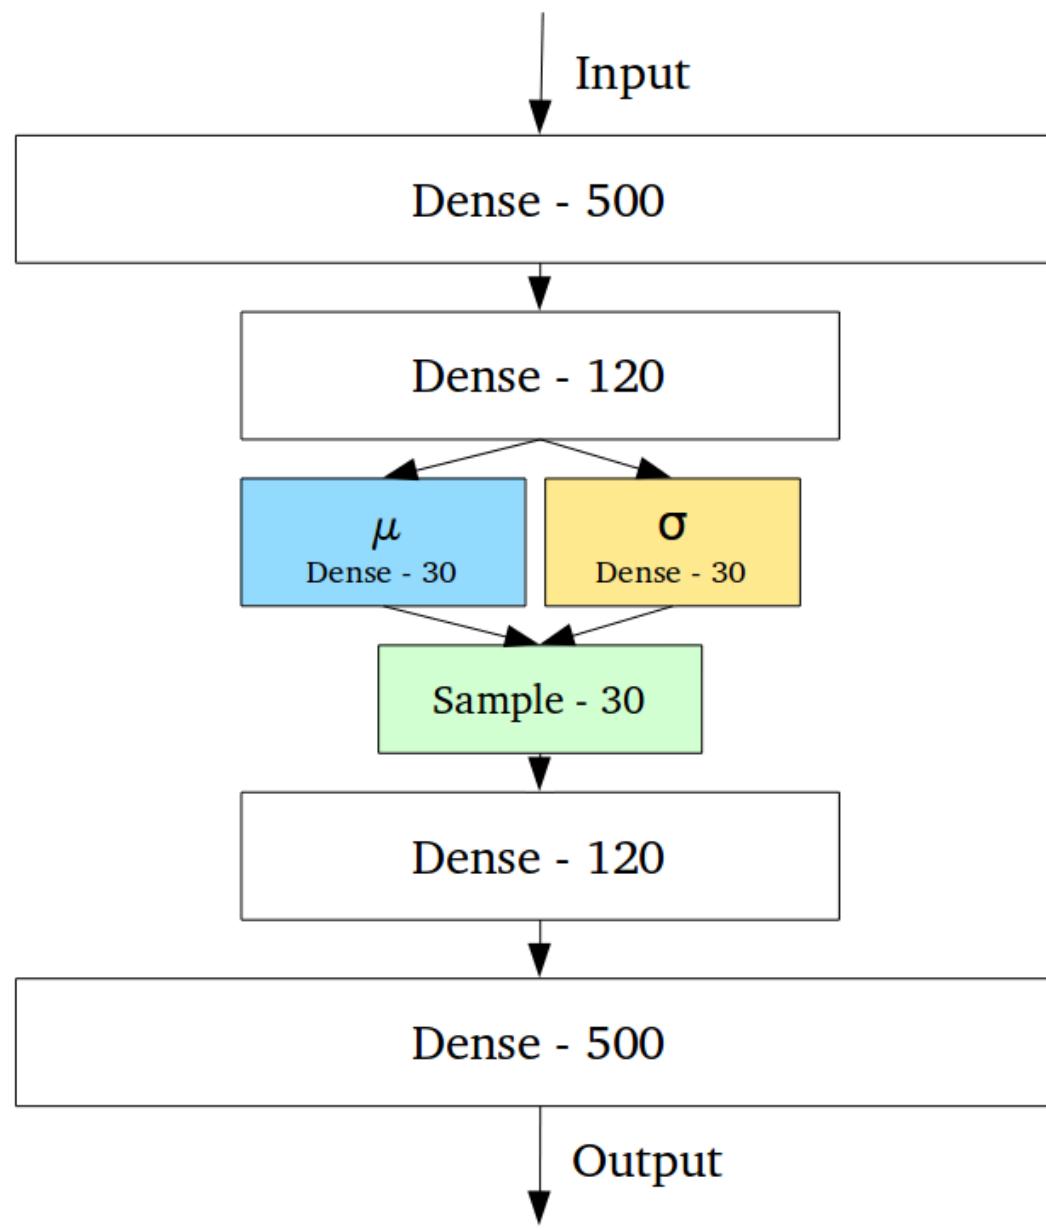
Variational Autoencoder (VAE)

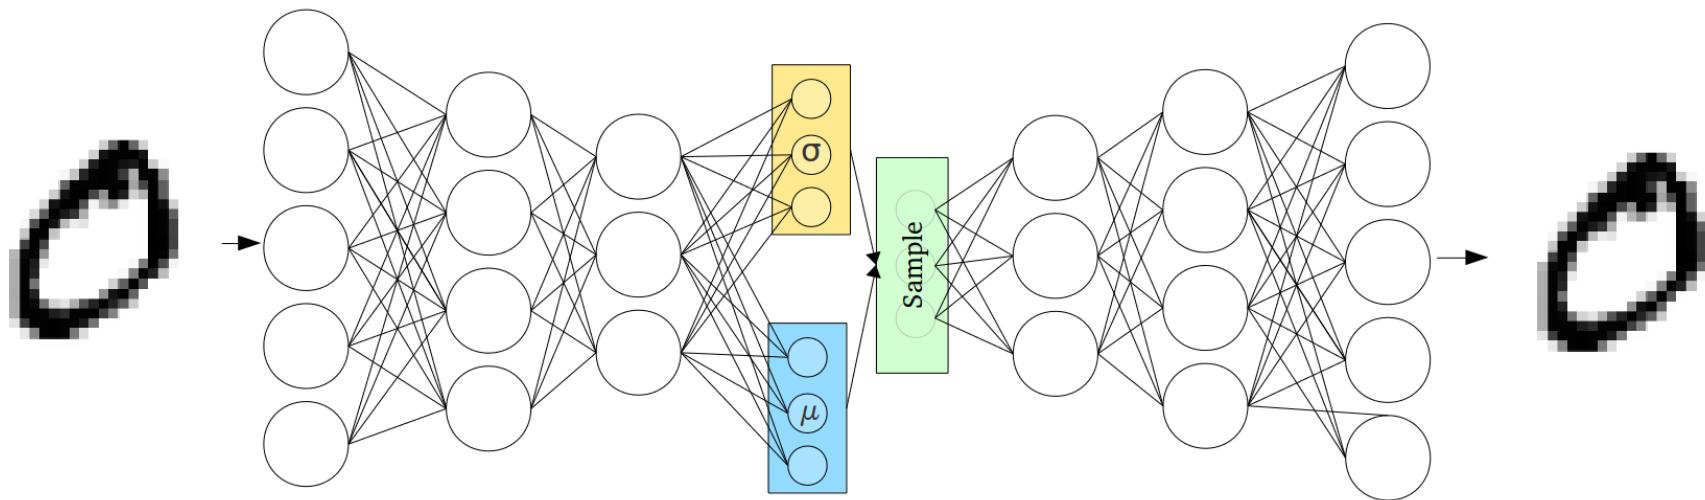


Standard Autoencoder
(direct encoding coordinates)



Variational Autoencoder
(μ and σ initialize a probability distribution)





Output
 μ

$[0.1, 1.2, 0.2, 0.8, \dots]$

Output
 σ

$[0.2, 0.5, 0.8, 1.3, \dots]$

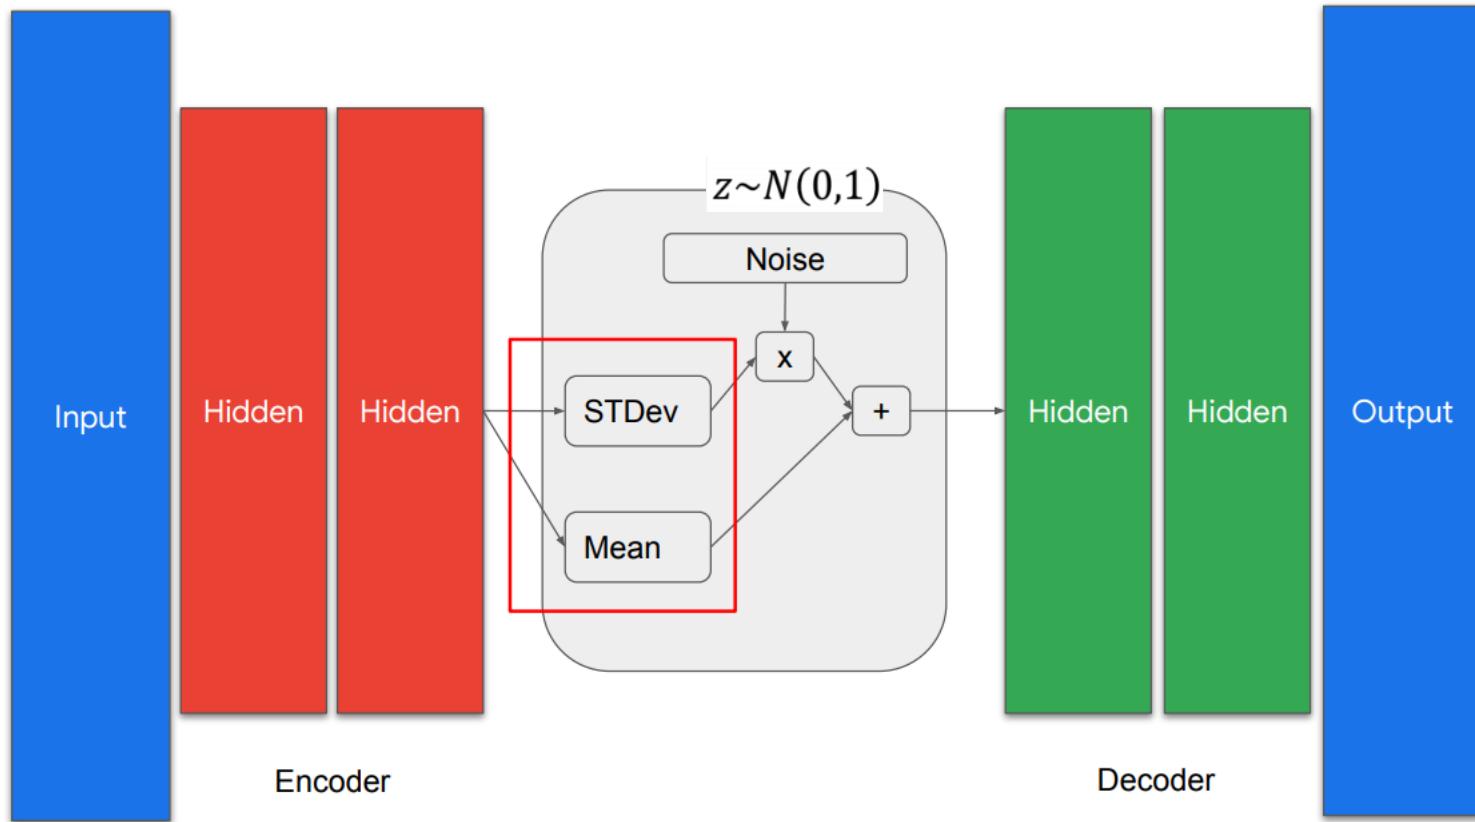
Intermediate
 X

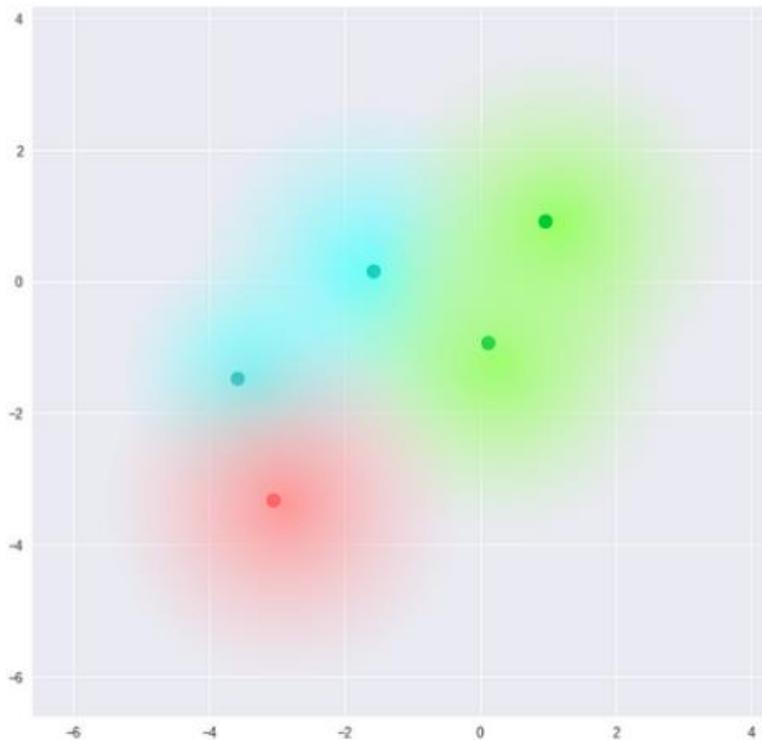
$[X_1 \sim N(0.1, 0.2^2), X_2 \sim N(1.2, 0.5^2), X_3 \sim N(0.2, 0.8^2), X_4 \sim N(0.8, 1.3^2), \dots]$

↓
sample

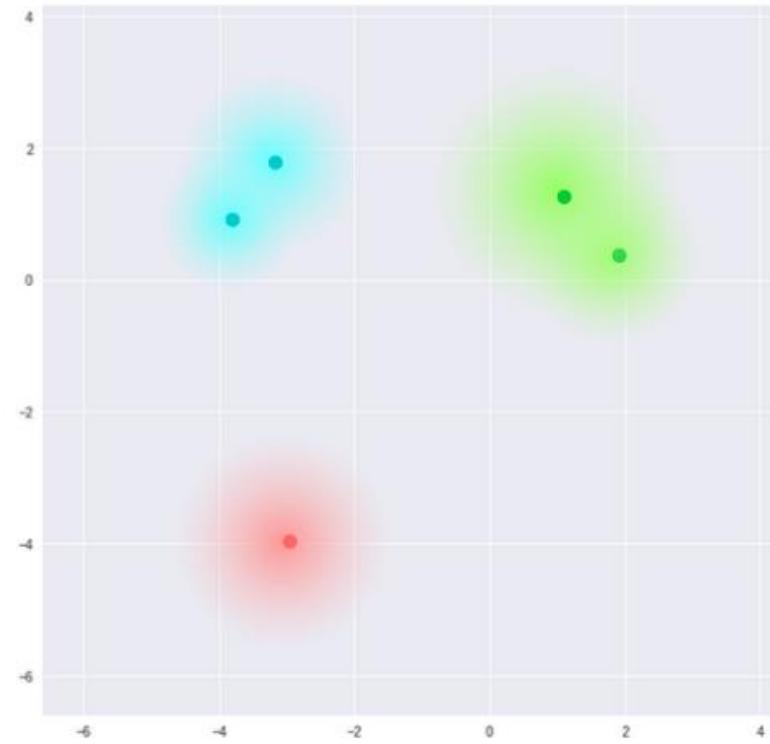
Sampled
vector

$[0.28, 1.65, 0.92, 1.98, \dots]$





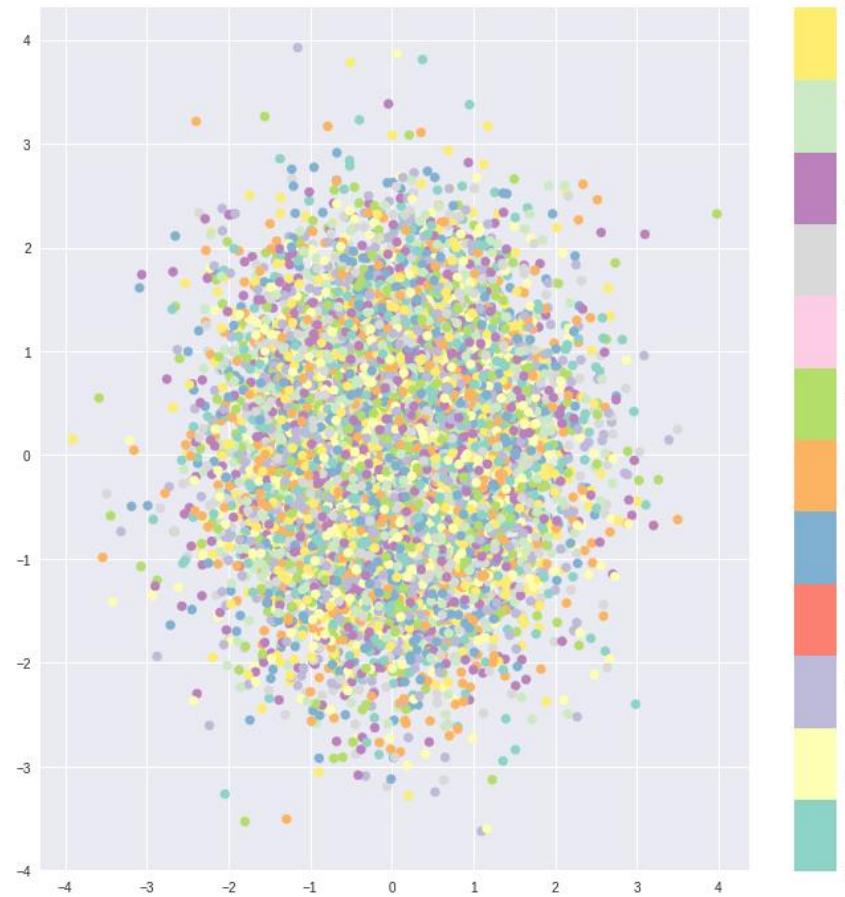
What we require

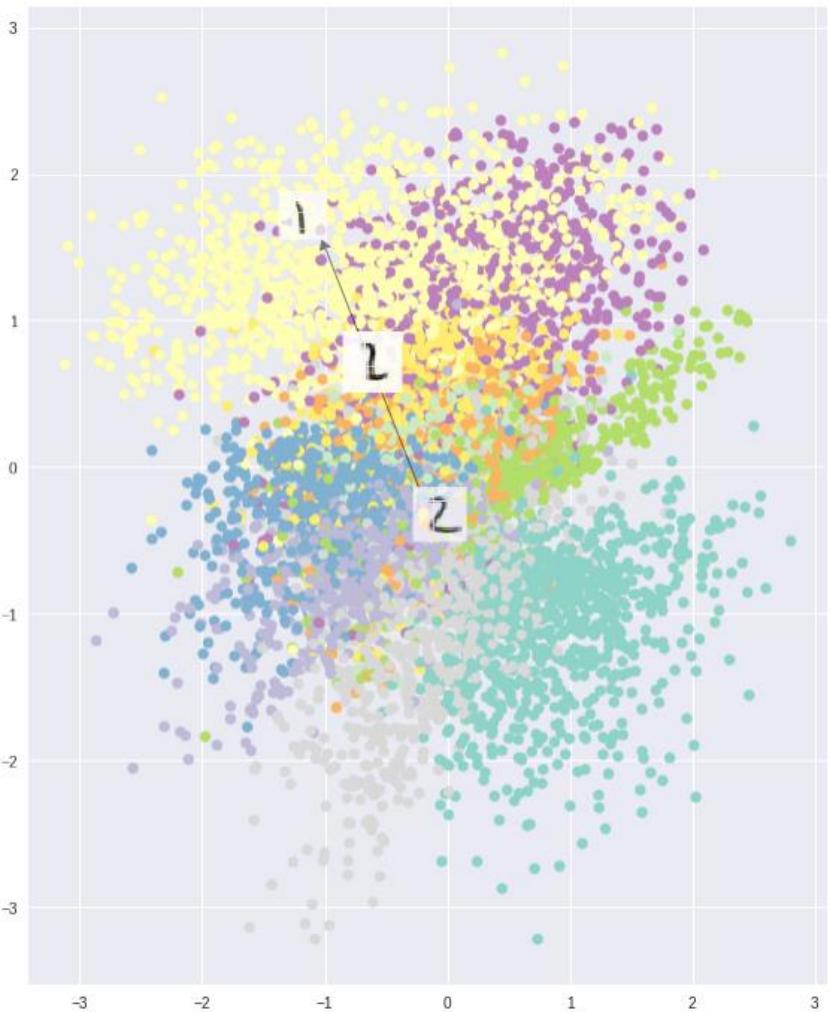


What we may inadvertently end up with

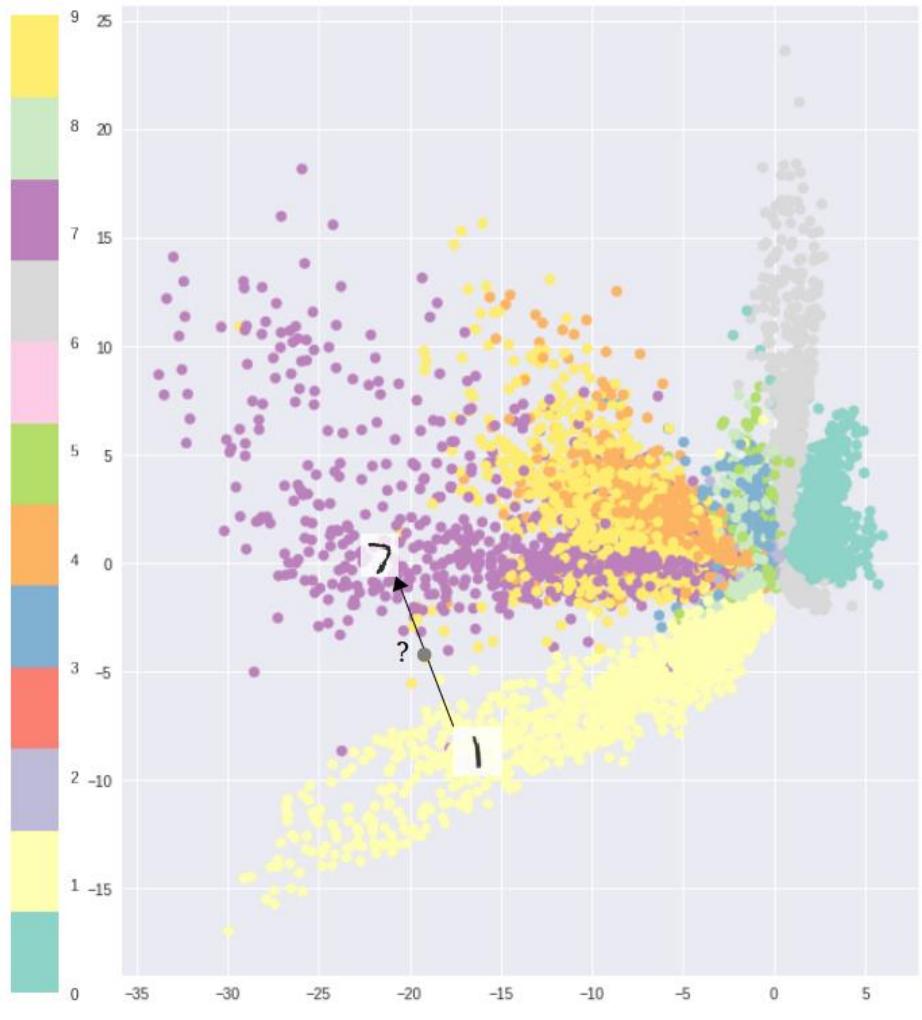
KL (Kullback–Leibler) divergence

$$\sum_{i=1}^n \sigma_i^2 + \mu_i^2 - \log(\sigma_i) - 1$$





VAE



AE

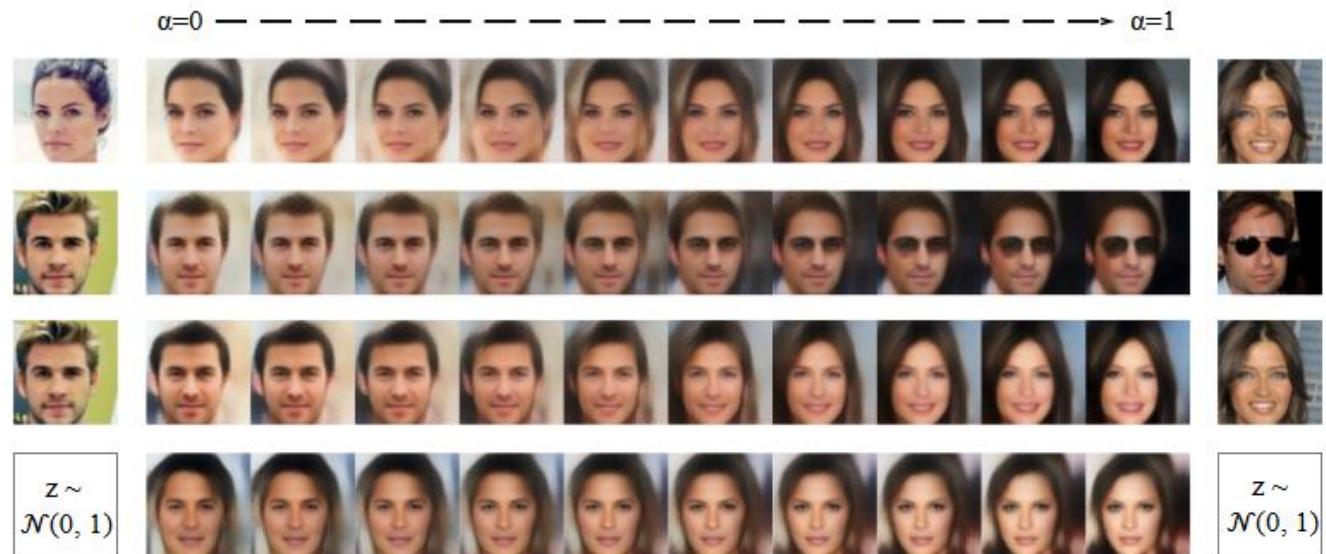


Figure 5. Linear interpolation for latent vector. Each row is the interpolation from left latent vector z_{left} to right latent vector z_{right} . e.g. $(1 - \alpha)z_{left} + \alpha z_{right}$. The first row is the transition from a non-smiling woman to a smiling woman, the second row is the transition from a man without eyeglass to a man with eyeglass, the third row is the transition from a man to a woman, and the last row is the transition between two fake faces decoded from $z \sim \mathcal{N}(0, 1)$.

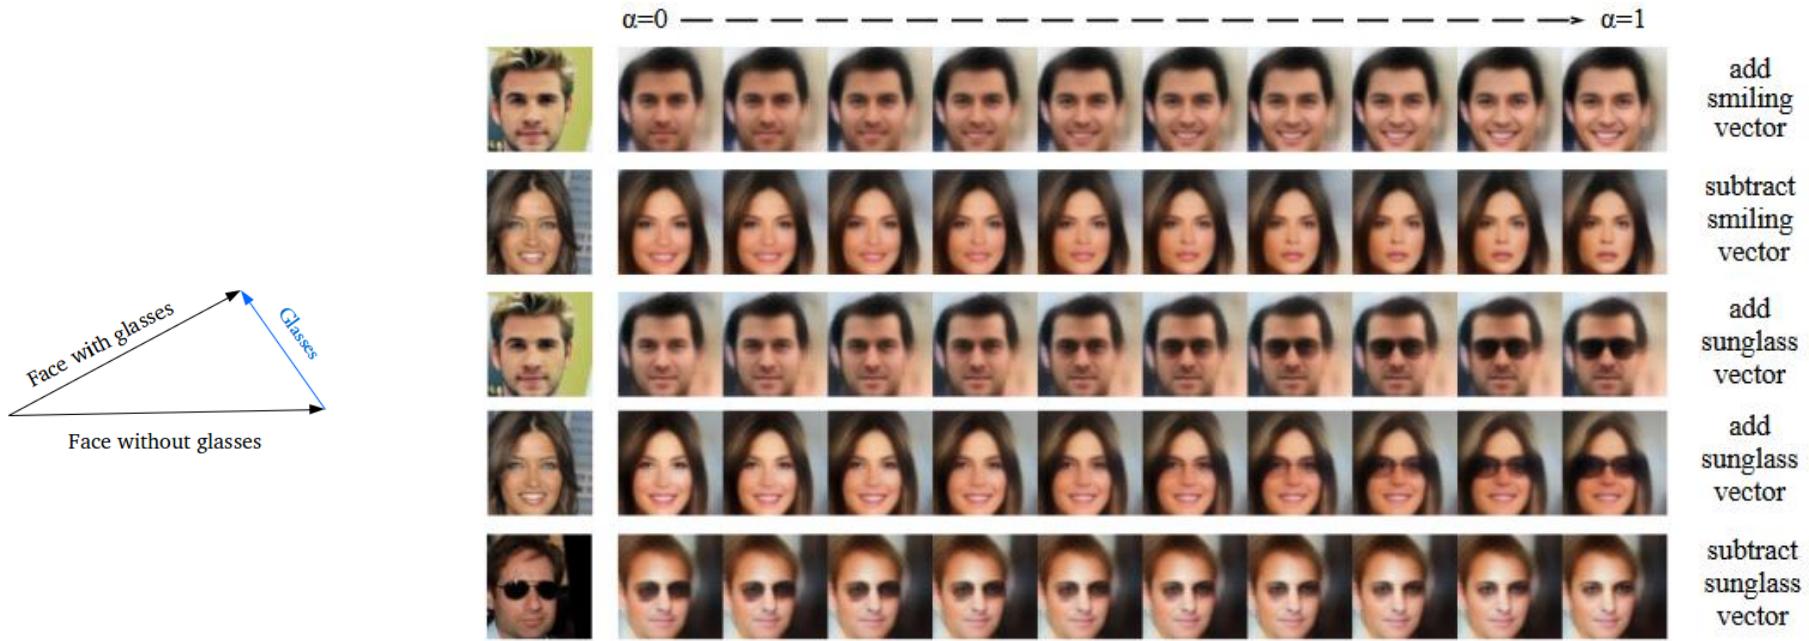
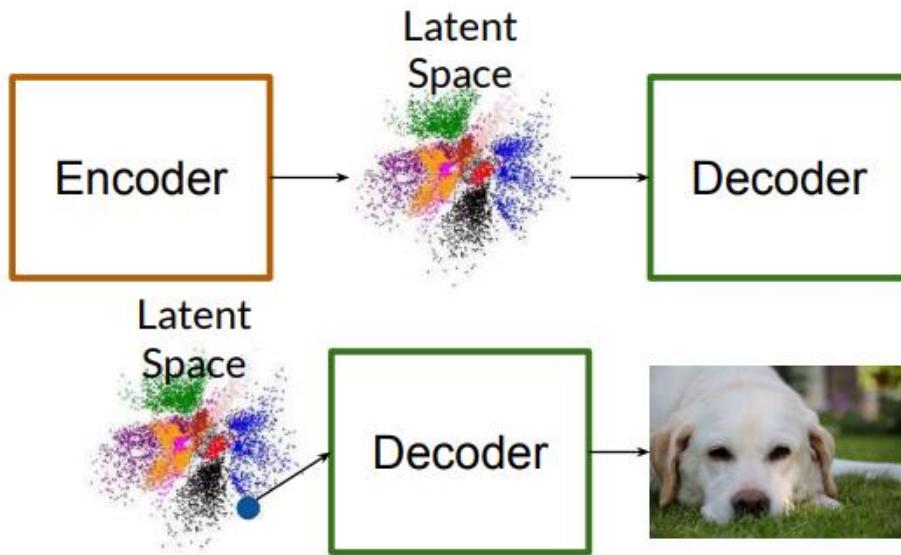


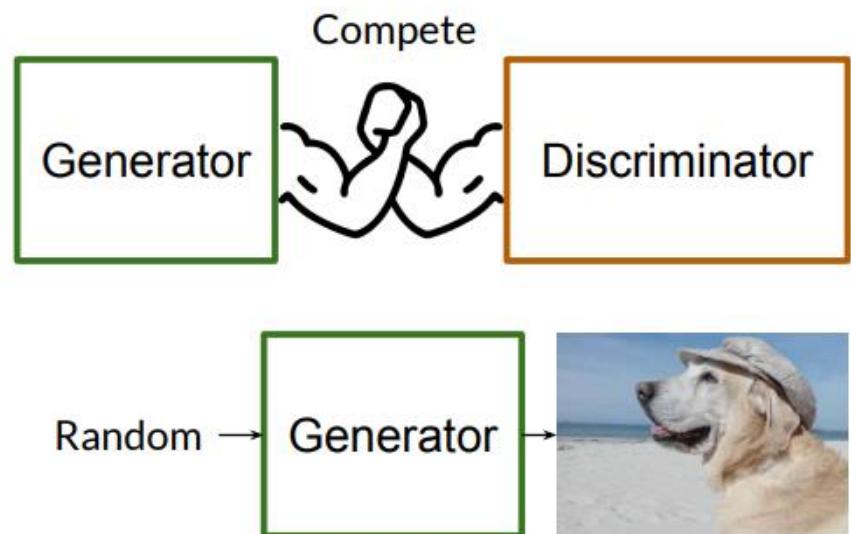
Figure 6. Vector arithmetic for visual attributes. Each row is the generated faces from latent vector z_{left} by adding or subtracting an attribute-specific vector, i.e., $z_{left} + \alpha z_{smiling}$, where $\alpha = 0, 0.1, \dots, 1$. The first row is the transition by adding a smiling vector with a linear factor α from left to right, the second row is the transition by subtracting a smiling vector, the third and fourth row are the results by adding a eyeglass vector to the latent representation for a man and women, and the last row shows results by subtracting an eyeglass vector.

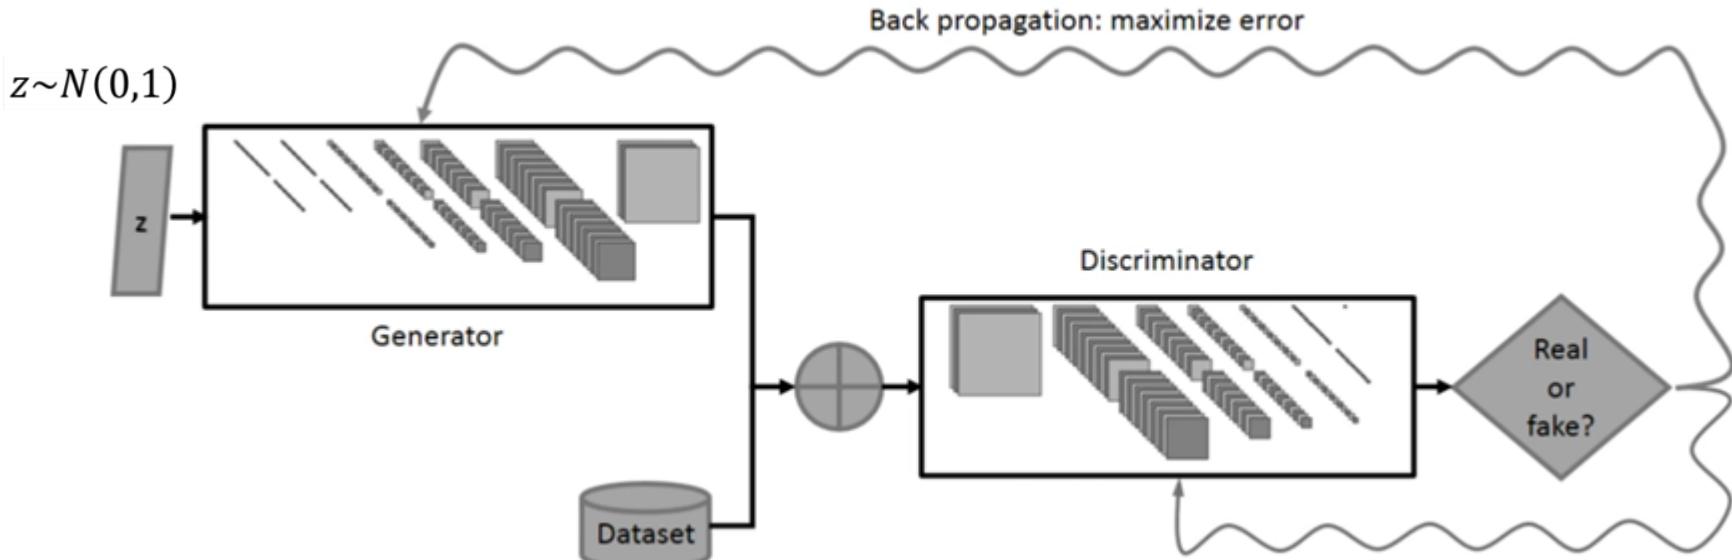
Generative Adversarial Networks (GAN)

Variational Autoencoders



Generative Adversarial Networks





```

for number of training iterations do
    for  $k$  steps do
        • Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
        • Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
        • Update the discriminator by ascending its stochastic gradient:
    
```

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

```

end for
    • Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
    • Update the generator by descending its stochastic gradient:

```

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

```
end for
```



a)



b)



c)



d)

Figure 2: Visualization of samples from the model. Rightmost column shows the nearest training example of the neighboring sample, in order to demonstrate that the model has not memorized the training set. Samples are fair random draws, not cherry-picked. Unlike most other visualizations of deep generative models, these images show actual samples from the model distributions, not conditional means given samples of hidden units. Moreover, these samples are uncorrelated because the sampling process does not depend on Markov chain mixing. a) MNIST b) TFD c) CIFAR-10 (fully connected model) d) CIFAR-10 (convolutional discriminator and “deconvolutional” generator)

“The GAN Zoo”

- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks

See also: <https://github.com/soumith/ganhacks> for tips and tricks for trainings GANs

- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

<https://github.com/hindupuravinash/the-gan-zoo>



Ian Goodfellow
@goodfellow_ian

4.5 years of GAN progress on face generation.

arxiv.org/abs/1406.2661 arxiv.org/abs/1511.06434

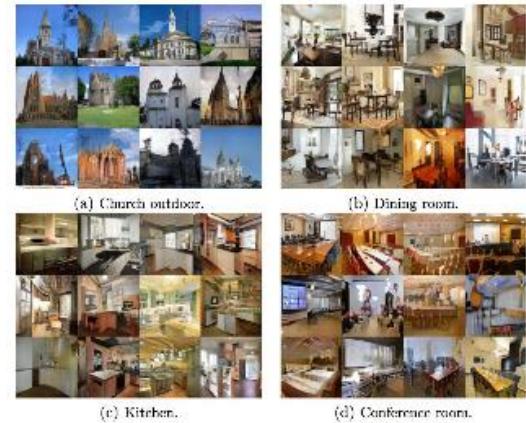
arxiv.org/abs/1606.07536 arxiv.org/abs/1710.10196

arxiv.org/abs/1812.04948



2017: Year of the GAN

Better training and generation

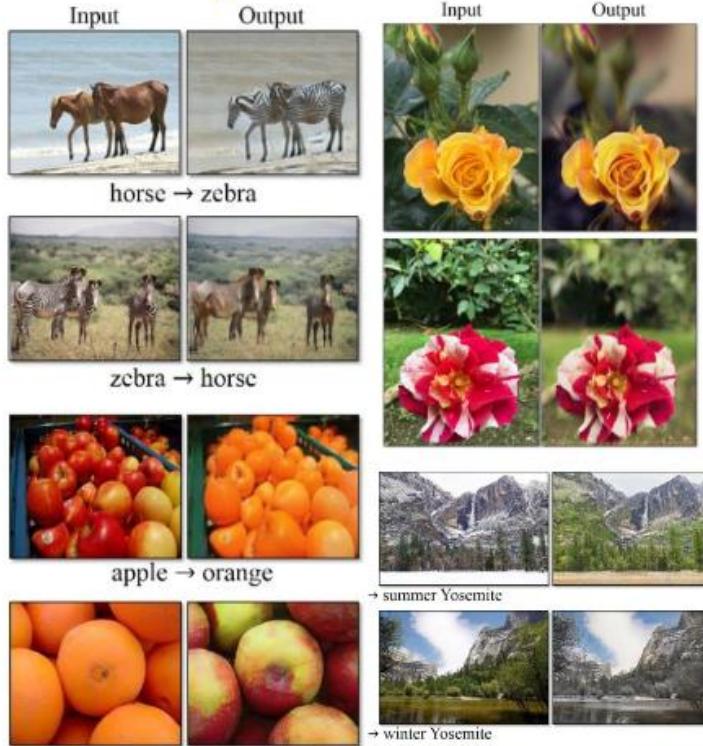


LSGAN. Mao et al. 2017.



BEGAN. Bertholet et al. 2017.

Source->Target domain transfer



CycleGAN. Zhu et al. 2017.

Text -> Image Synthesis

this small bird has a pink breast and crown, and black primaries and secondaries.



Reed et al. 2017.

this magnificent fellow is almost all black with a red crest, and white cheek patch.



Many GAN applications



Pix2pix. Isola 2017. Many examples at <https://phillipi.github.io/pix2pix/>



Mao et al. (2016b) (128 × 128)

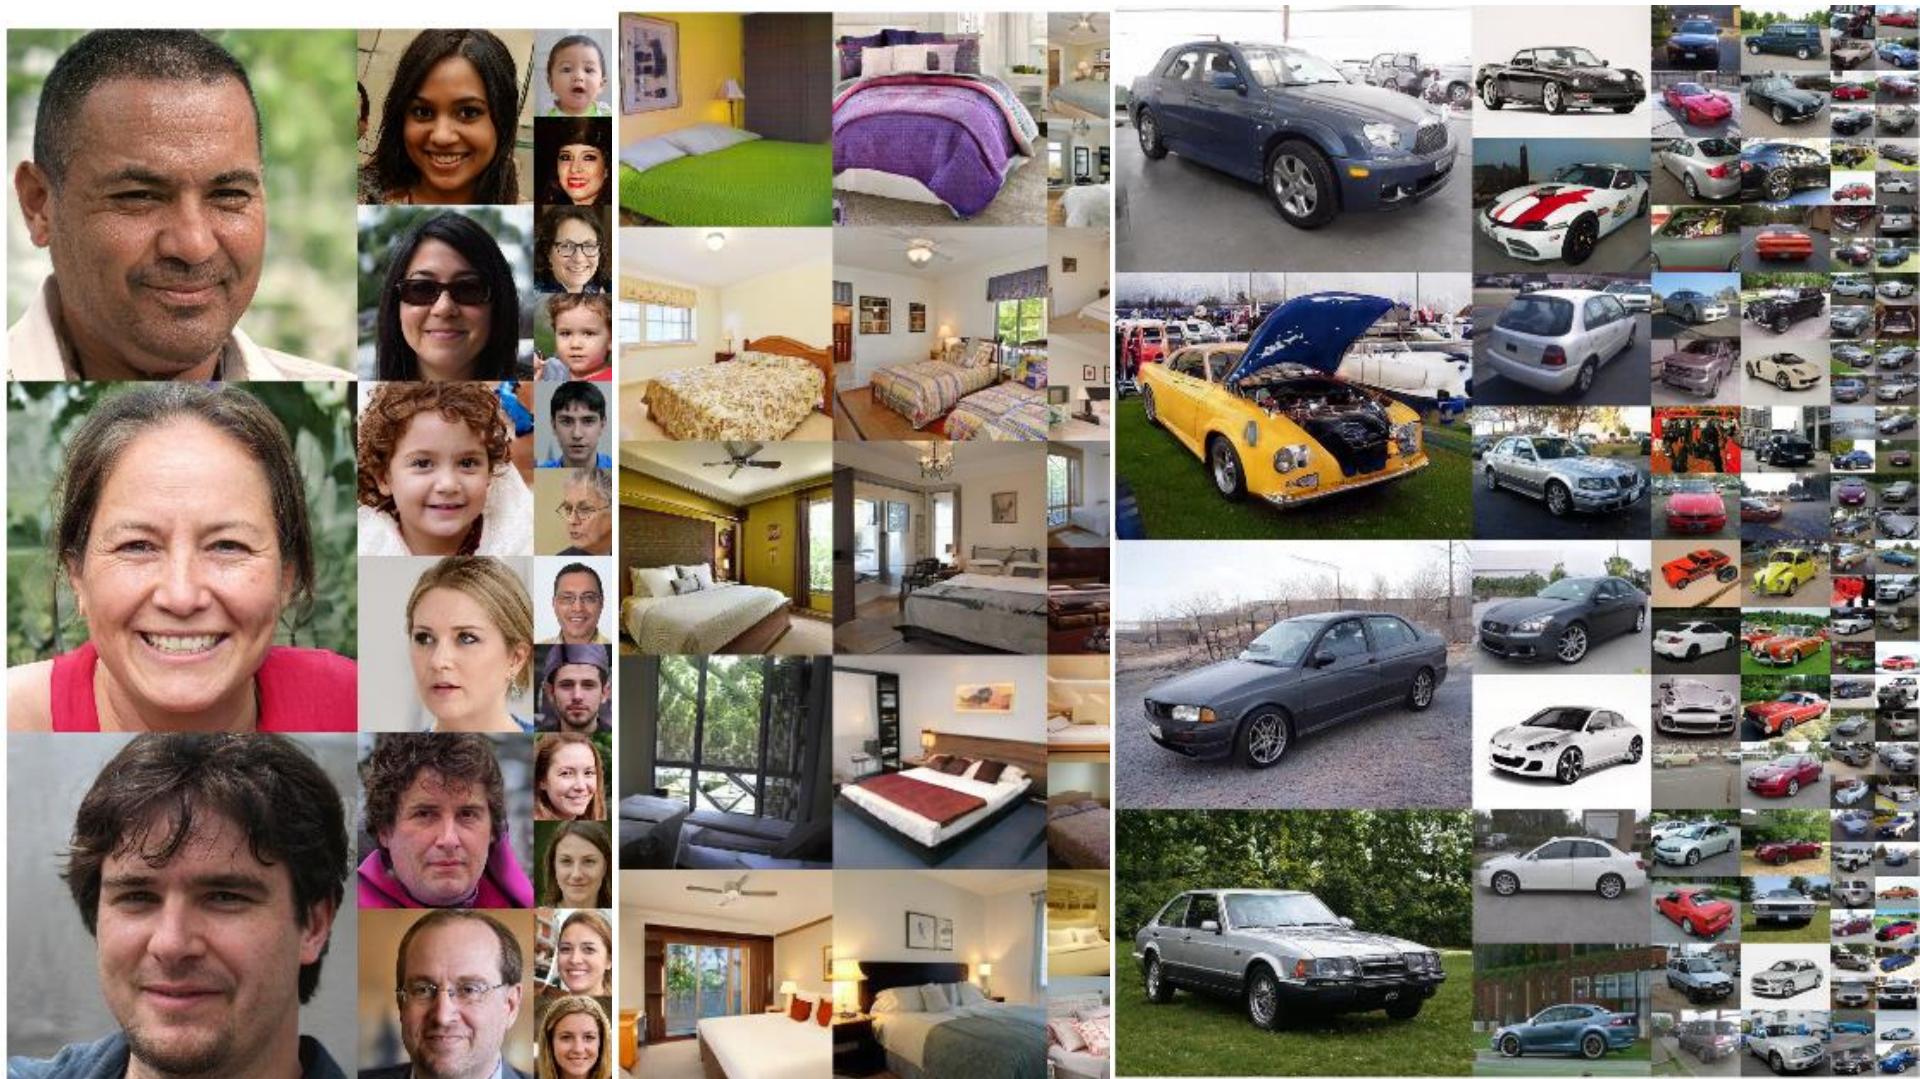
Gulrajani et al. (2017) (128 × 128)

Our (256 × 256)

Figure 6: Visual quality comparison in LSUN BEDROOM; pictures copied from the cited articles.



StyleGAN'18



<https://thispersondoesnotexist.com/>



Imagined by a GAN (generative adversarial network)
StyleGAN2 (Dec 2019) - Karras et al. and Nvidia
Don't panic, it's AI art [works] [1] [2] [3]
About AI continue to dream | Contact me
Code for training your own [original] [simple]
Art • Cats • Horses • Chemicals
Another

<https://generated.photos/faces>

A screenshot of the Generated Photos website. The interface includes a navigation bar with links for Faces, Use cases, Datasets, API, Pricing, and Sign In. On the left, there are several filter panels: "Select Photos" (radio buttons for All 2,687,765 with current filter, Random 100, or All 30 on this page), "Background Color" (dropdown menu), "Face" (radio buttons for All, Natural NEW, or Beautified), "Head Pose", "Sex", "Age", "Ethnicity", "Eye Color", "Hair Color", "Hair Length", and "Emotion". To the right, a grid of diverse faces is displayed in a 4x6 layout.

Text-to-Image Generation

Text
description



64x64
GAN-INT-CLS

128x128
GAWWN

256x256
StackGAN-v1

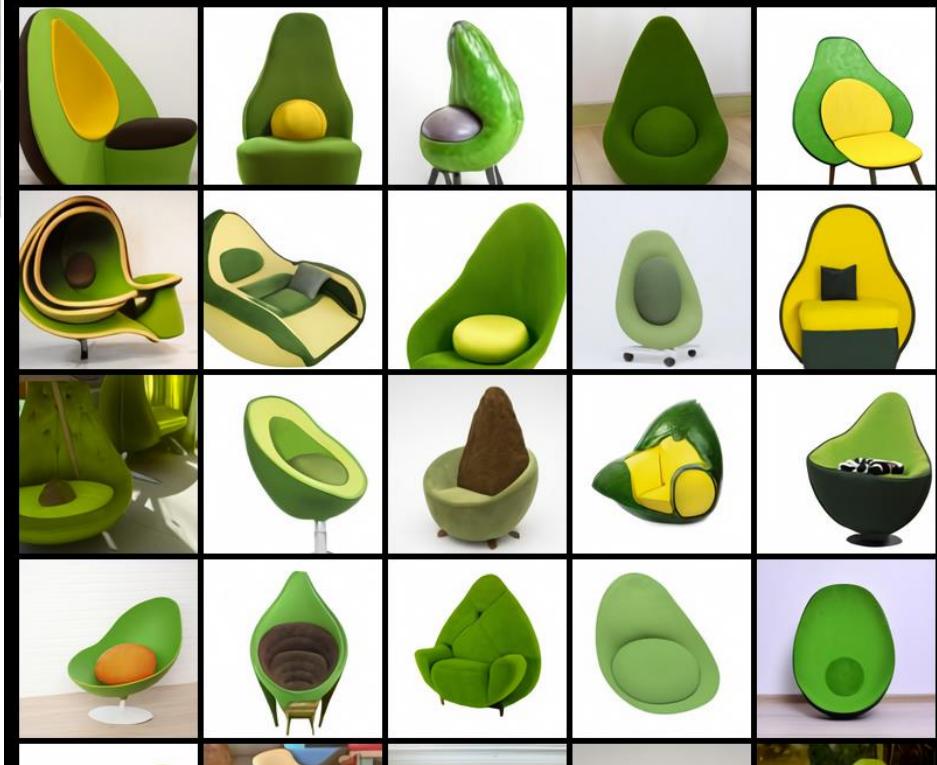
256x256
StackGAN-v2

<https://paperswithcode.com/task/text-to-image-generation>

TEXT PROMPT

an armchair in the shape of an avocado. an armchair imitating an avocado.

AI-GENERATED IMAGES



<https://openai.com/blog/dall-e/>

The image shows the 'Community Showcase' page of the stability.ai website. At the top right, there's a large title 'Community Showcase'. Below it is a grid of various AI-generated images, including a blue-skinned woman with intricate patterns, a white owl, a portrait of Karl Marx, a butterfly-shaped necklace, a woman in a golden armor, and several smaller images of landscapes and people. On the left side, there's a sidebar with sections like 'Join the Beta', 'Stability', 'Showcase', 'Documentation', and 'Sign In'. The bottom navigation bar includes 'Models', 'Applications', and 'Deploy'.

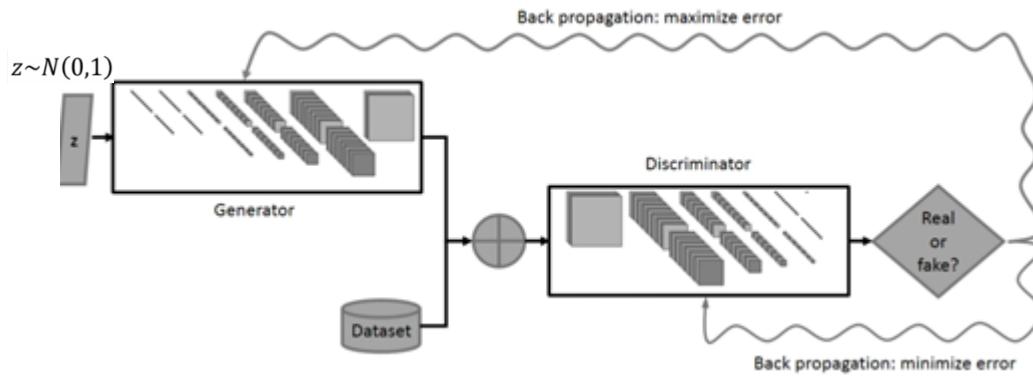
The image shows the homepage of Starfee.ai. In the top left corner is the company logo, which consists of a stylized blue and green flower-like icon followed by the text "Starfee.ai". To the right of the logo is a horizontal row of five small images generated by the AI: a person's face, a hamster in a red dragon costume, a woman in a yellow floral headdress, a cat looking at a screen, and a mountain landscape. To the far right of these images is a large, bold number "36,324" and the text "Кількість створених робіт" (Number of created works). Below the main navigation bar, there is a section with a button labeled "+143 заготовлені стилі" (143 prepared styles) accompanied by three small circular icons. The central part of the page features a large, bold headline: "Найкраща Платформа для створення Логотипів, Креативів" (The best platform for creating logos, creative works). At the bottom of the page, there are two small images: one of a person in a desert and another of a colorful lizard. On the right side, there is a box containing the text "Роботи згенеровані за допомогою" (Works generated with the help of).

GANs for 3D Objects



Wu, Jiajun, et al. "Learning a probabilistic latent space of object shapes via 3d generative-adversarial modeling." *Advances in neural information processing systems*. 2016.

Unconditional



Examples from *random classes*

Training dataset *doesn't need to be labeled*

Controllable

Examples with the **features that you want**

Training dataset **doesn't need to be labeled**

Manipulate the z vector input

Conditional

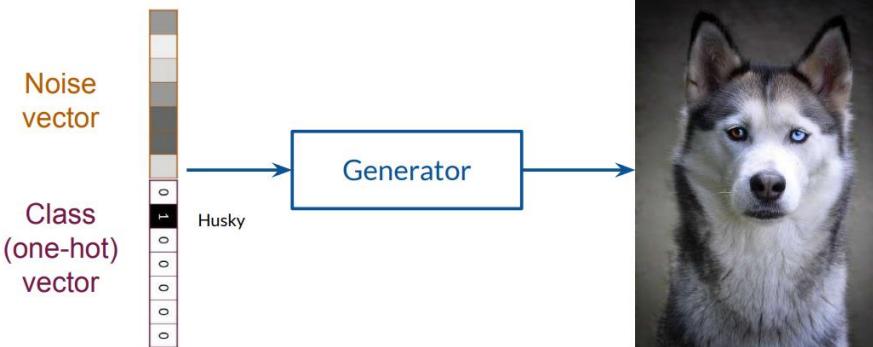
Examples from *the classes you want*

Training dataset *needs to be labeled*

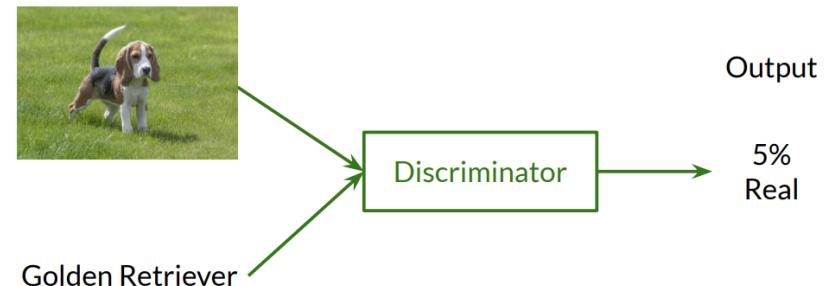
Append a class vector to the input

Conditional generation

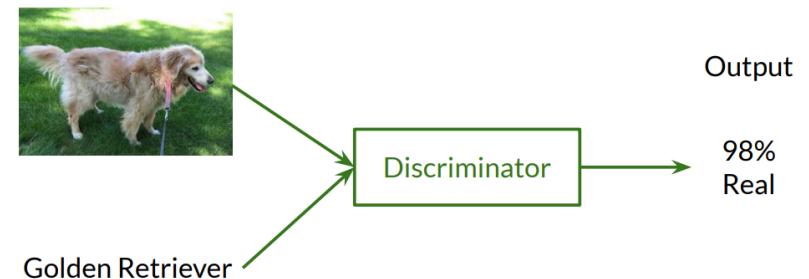
Generator Input



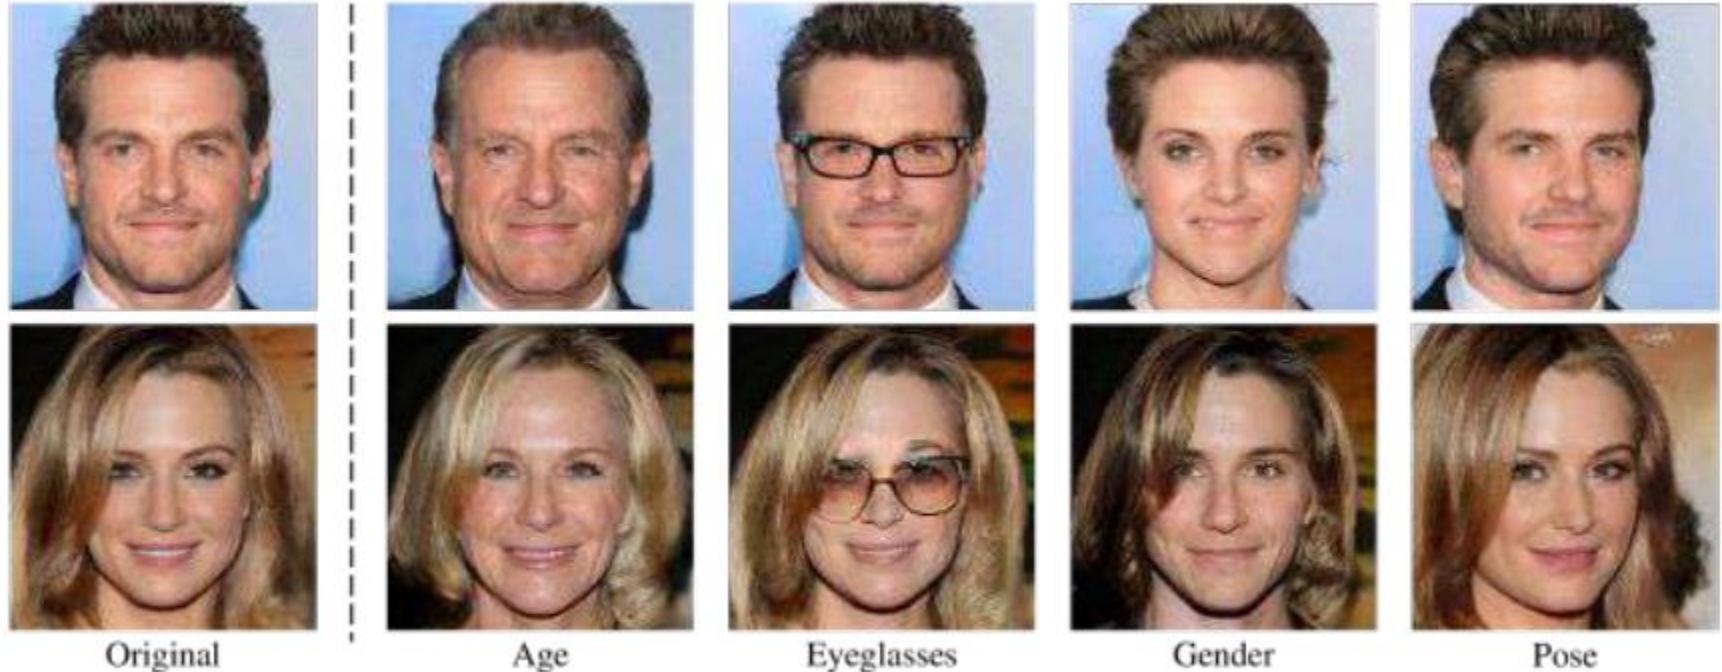
Discriminator Input



Discriminator Input



Controllable generation

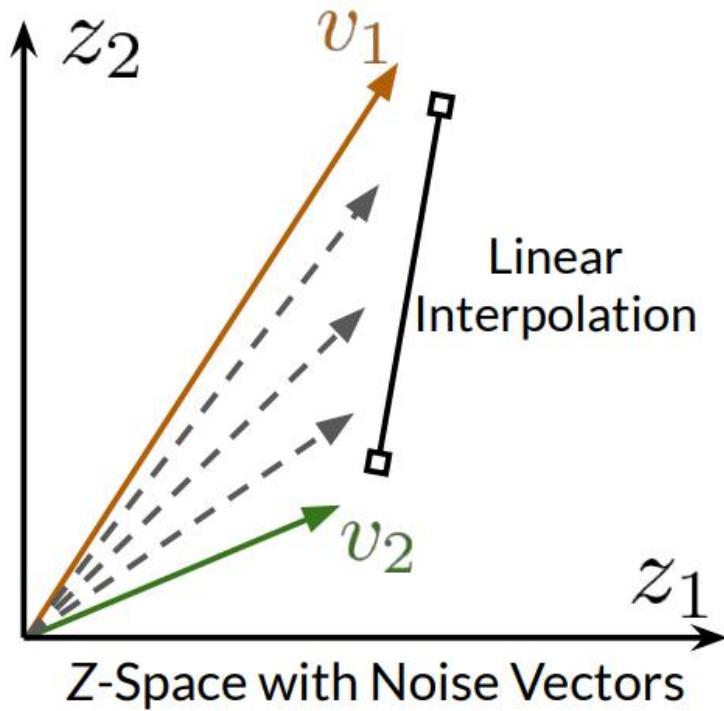


Change specific features of the output

Tweak the input noise vector to get different features on the output

[1907.10786.pdf \(arxiv.org\)](https://arxiv.org/pdf/1907.10786.pdf)

Interpolation Using the Z-Space



Adversarial attack



x
“panda”
57.7% confidence

$+ .007 \times$



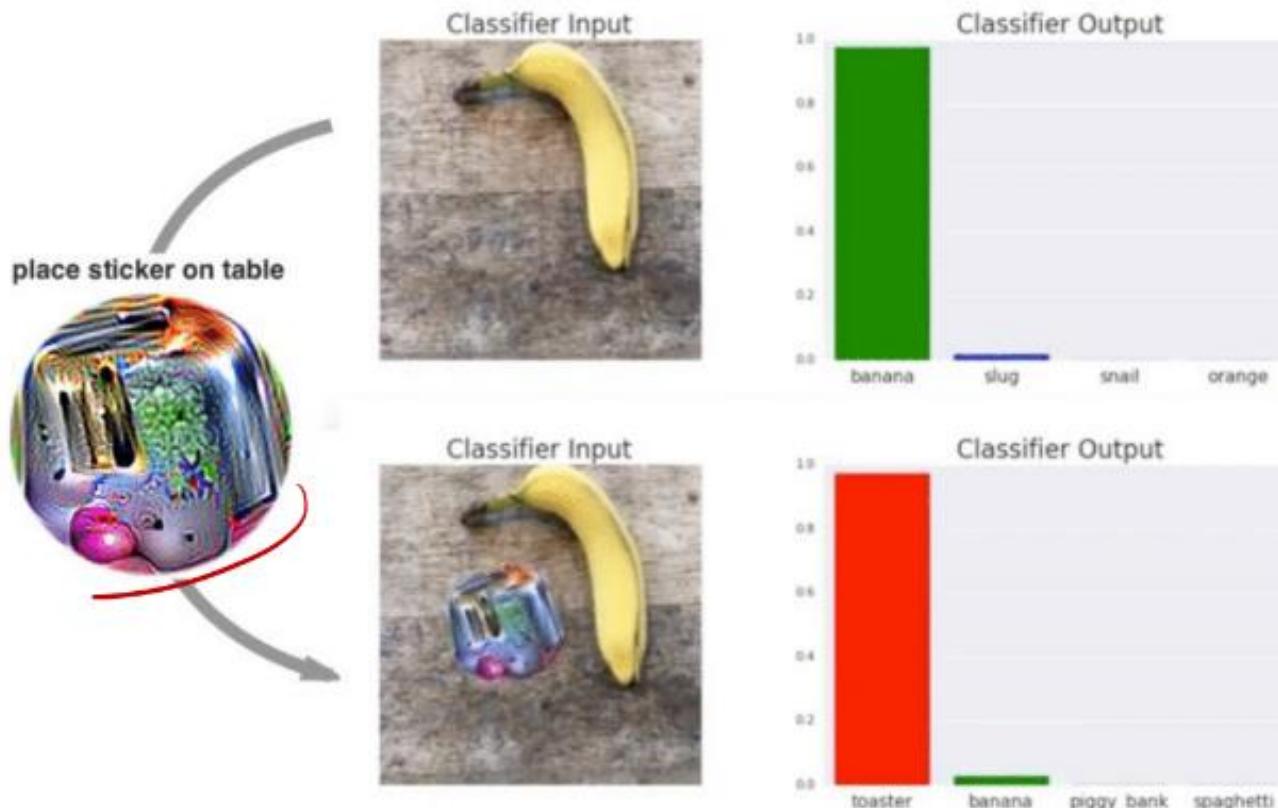
$\text{sign}(\nabla_x J(\theta, x, y))$
“nematode”
8.2% confidence

=



$x +$
 $\epsilon \text{sign}(\nabla_x J(\theta, x, y))$
“gibbon”
99.3 % confidence

Adversarial Patch





Convolutional Neural Networks

DeepLearning.AI

Convolutional Neural Networks

by DeepLearning.AI

курс: загальні відомості

In the fourth course of the Deep Learning Specialization, you will understand how computer vision has evolved and become familiar with its exciting applications such as autonomous driving, face recognition, reading radiology images, and more.

▽ Матеріал курсу

- Тиждень 1
- Тиждень 2
- Тиждень 3
- Тиждень 4

By the end, you will be able to build a convolutional neural network, including recent variations such as residual networks; apply convolutional networks to visual detection and recognition tasks; and use neural style transfer to generate art and apply these algorithms to a variety of image, video, and other 2D or 3D data.

The Deep Learning Specialization is our foundational program that will help you understand the capabilities, challenges, and consequences of deep learning and prepare you to participate in the development of leading-edge AI technology. It provides a pathway for you to gain the knowledge and skills to apply machine learning to your work, level up your technical career, and take the definitive step in the world of AI.

 Показати менше

Оцінки

Примітки

Повідомлення

Ресурси



Викладають: [Andrew Ng](#), Instructor

Founder, DeepLearning.AI & Co-founder, Coursera



Advanced Computer Vision with
TensorFlow
DeepLearning.AI

Advanced Computer Vision with TensorFlow

by DeepLearning.AI

Матеріал курсу

- Тиждень 1
- Тиждень 2
- Тиждень 3
- Тиждень 4

курс: загальні відомості

In this course, you will:

- a) Explore image classification, image segmentation, object localization, and object detection. Apply transfer learning to object localization and detection.
- b) Apply object detection models such as regional-CNN and ResNet-50, customize existing models, and build your own models to detect, localize, and label your own rubber duck images.
- c) Implement image segmentation using variations of the fully convolutional network (FCN) including U-Net and d) Mask-RCNN to identify and detect numbers, pets, zombies, and more.
- d) Identify which parts of an image are being used by your model to make its predictions using class activation maps and saliency maps and apply these ML interpretation methods to inspect and improve the design of a famous network, AlexNet.

The DeepLearning.AI TensorFlow: Advanced Techniques Specialization introduces the features of TensorFlow that provide learners with more control over their model architecture and tools that help them create and train advanced ML models.

This Specialization is for early and mid-career software and machine learning engineers with a foundational understanding of TensorFlow who are looking to expand their knowledge and skill set by learning advanced TensorFlow features to build powerful models.

[▲ Показати менше](#)

Оцінки

Примітки

Повідомлення

Інформація



Викладають: [Laurence Moroney](#), Instructor

Lead AI Advocate, Google



Generative Deep Learning with TensorFlow

DeepLearning.AI

Generative Deep Learning with TensorFlow

by DeepLearning.AI

курс: загальні відомості

In this course, you will:

- a) Learn neural style transfer using transfer learning: extract the content of an image (eg. swan), and the style of a painting (eg. cubist or impressionist), and combine the content and style into a new image.
- b) Build simple AutoEncoders on the familiar MNIST dataset, and more complex deep and convolutional architectures on the Fashion MNIST dataset, understand the difference in results of the DNN and CNN AutoEncoder models, identify ways to de-noise noisy images, and build a CNN AutoEncoder using TensorFlow to output a clean image from a noisy one.
- c) Explore Variational AutoEncoders (VAEs) to generate entirely new data, and generate anime faces to compare them against reference images.
- d) Learn about GANs; their invention, properties, architecture, and how they vary from VAEs, understand the function of the generator and the discriminator within the model, the concept of 2 training phases and the role of introduced noise, and build your own GAN that can generate faces.

The DeepLearning.AI TensorFlow: Advanced Techniques Specialization introduces the features of TensorFlow that provide learners with more control over their model architecture, and gives them the tools to create and train advanced ML models.

This Specialization is for early and mid-career software and machine learning engineers with a foundational understanding of TensorFlow who are looking to expand their knowledge and skill set by learning advanced TensorFlow features to build powerful models.

[▲ Показати менше](#)

Оцінки

Примітки

Повідомлення 1

Інформація



Викладають: [Laurence Moroney](#), Instructor

Lead AI Advocate, Google

Examples

▶ Segmentation

https://colab.research.google.com/github/https-deeplearning-ai/tensorflow-3-public/blob/main/Course%203%20-%20Advance%20Computer%20Vision/W3/ungraded_labs/C3_W3_Lab_2_OxfordPets-UNet.ipynb#scrollTo=_xzl28AfxFQi

▶ Object_Localization

- ▶ https://colab.research.google.com/github/https-deeplearning-ai/tensorflow-3-public/blob/main/Course%203%20-%20Advance%20Computer%20Vision/W1/ungraded_labs/C3_W1_Lab_3_Object_Localization.ipynb

▶ Object Detection with TF Hub

- ▶ https://colab.research.google.com/github/https-deeplearning-ai/tensorflow-3-public/blob/main/Course%203%20-%20Advance%20Computer%20Vision/W2/ungraded_labs/C3_W2_Lab_1_Simple_Object_Detection.ipynb
- ▶ https://colab.research.google.com/github/https-deeplearning-ai/tensorflow-3-public/blob/main/Course%203%20-%20Advance%20Computer%20Vision/W2/ungraded_labs/C3_W2_Lab_2_Object_Detection.ipynb#scrollTo=CxmDMK4yupqg
- ▶ https://colab.research.google.com/github/tensorflow/hub/blob/master/examples/colab/tf2_object_detection.ipynb
- ▶ https://colab.research.google.com/github/tensorflow/models/blob/master/research/object_detection/colab_tutorials/eager_few_shot_od_training_tf2_colab.ipynb

First GAN

- ▶ https://colab.research.google.com/github/https-deeplearning-ai/tensorflow-3-public/blob/main/Course%204%20-%20Generative%20Deep%20Learning/W4/ungraded_labs/C4_W4_Lab_1_First_GAN.ipynb